

* * * * * With Best-Selling Instructors Chris Dutton & Aaron Parry



COURSE STRUCTURE

This is a **project-based course** designed for students looking for a practical, hands-on, and highly engaging approach to learning Power BI Desktop for business intelligence

Course resources include:



Downloadable PDF eBook (200+ pages) containing all course slides, assignments and reference materials



Quizzes and **Assignments** to reinforce key concepts and simulate real-world scenarios, with step-by-step solution videos



Complete **Bonus Project** to test your abilities and apply the skills developed throughout the course to a brand-new data set

COURSE OUTLINE

1 Introducing Power BI Desktop	Installing Power BI Desktop, exploring the Power BI workflow, comparing Power BI vs. Excel, etc.
Connecting & Shaping Data	Connecting to data, shaping & transforming tables, using profiling tools, editing, merging & appending queries, etc.
3 Creating a Data Model	Building relational models, creating table relationships, understanding cardinality and filter flow, etc.
Calculating Measures with DAX	Understanding DAX syntax, adding calculated columns and measures, writing common formulas and functions, etc.
Visualizing Data with Dashboards	Inserting charts and visuals, customizing formats, editing interactions, applying filters and bookmarks, etc.
Optimizing Power BI Performance	Exploring common Power BI optimization tools within the Optimize and External tools menus

COURSE PROJECT



You've just been hired as a Business Intelligence Analyst by **AdventureWorks***, a global manufacturing company that produces cycling equipment and accessories



The management team needs a way to track KPIs (sales, revenue, profit, returns), compare regional performance, analyze product-level trends, and identify high-value customers.

All you've been given is a **folder of raw csv files**, which contain information about transactions, returns, products, customers, and sales territories.



Use Power BI Desktop to:

- Connect and transform the raw data
- Build a relational data model
- Create calculated columns and measures with DAX
- Design an interactive dashboard to visualize the data



SETTING EXPECTATIONS

- What you see on your screen may not always match mine
 - Power BI Desktop features are updated frequently, with new versions released each month
 - **NOTE:** Power BI is currently only compatible with PC/Windows (not available for Mac)
- This course is designed to help you build foundational skills
 - Our goal is to help you build a deep foundational understanding of the Power BI desktop workflow; some topics may be simplified, and we won't cover some advanced tools (M code, advanced DAX, R/Python visuals, etc.)
- This is a **hands-on** and **project-based** learning experience
 - You will get the most value out of this course if you follow along closely with the demos and assignments; we'll be working through the entire BI workflow to create a professional-quality dashboard from scratch
- We will not cover **Power BI Service** as part of this course
 - This course focuses on Power BI Desktop specifically; online sharing and collaboration features (app.powerbi.com) require a separate account and are covered in-depth in a separate course



MEET POWER BI



In this section we'll **introduce Power BI Desktop**, review the download and installation process, adjust default settings, and explore the Power BI interface and workflow

TOPICS WE'LL COVER:

Introducing Power BI

Power BI vs. Excel

Installation Options

Adjusting Settings

Interface & Workflow

Helpful Resources

GOALS FOR THIS SECTION:

- Download and install Power BI Desktop, and adjust the settings for our course project
- Understand the role that Power BI plays within the broader Microsoft ecosystem
- Explore core components of the Power BI Desktop interface
- Review the business intelligence workflow that we'll follow as we build our course project

MEET POWER BI



Microsoft Power BI is a self-service business intelligence platform, which includes both desktop and web-based applications for connecting, modeling, and visualizing data

Learn more at **powerbi.microsoft.com**



Figure 1: Magic Quadrant for Analytics and Business Intelligence Platforms



WHY POWER BI?



Connect, transform and load millions of rows of data

• Access data from virtually anywhere (database tables, flat files, web, cloud services, folders, etc.), and create fully automated workflows to extract, transform and load data for analysis



Build relational models to blend data from multiple sources

Create table relationships to analyze holistic performance across an entire relational data model



Define complex calculations using Data Analysis Expressions (DAX)

Enhance datasets and enable advanced analytics with powerful and portable DAX expressions



Bring data to life with interactive reports and dashboards

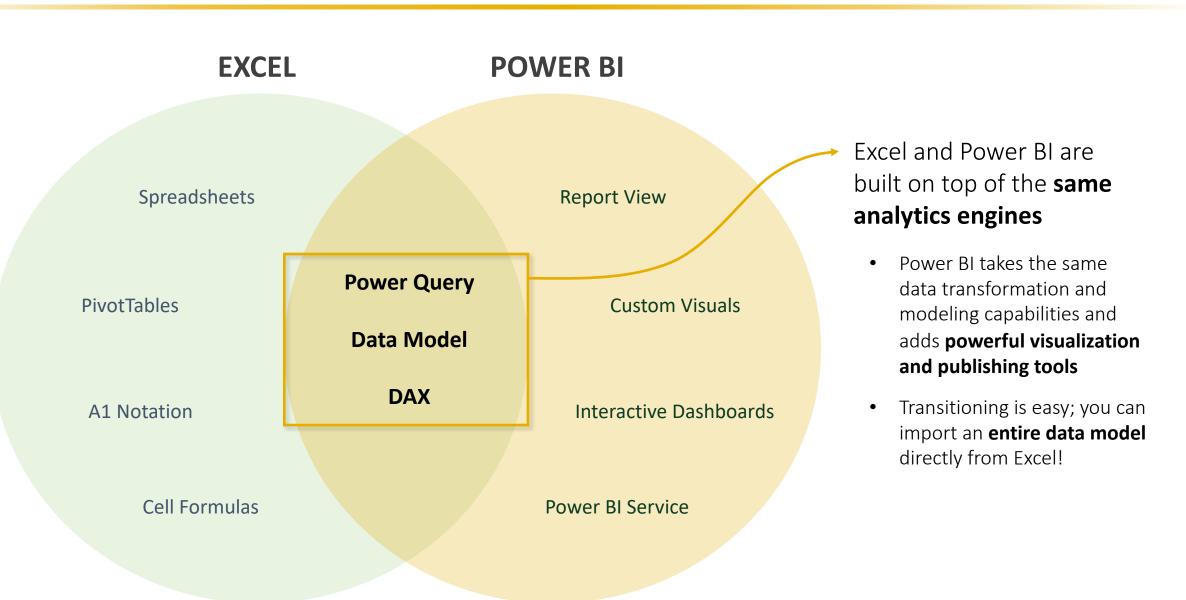
• Build professional-quality reports and dashboards with best-in-class visualization tools



Develop a versatile, in-demand skill set

Power BI is the industry leader in self-service BI, and the skills you build in this course will be highly transferrable

EXCEL VS. POWER BI



INSTALLING POWER BI DESKTOP

1) Download from Microsoft store apps.microsoft.com



- Windows handles automatic updates
- Updates only elements that have been changed
- Doesn't require administrator access

2) Download manually from web powerbi.microsoft.com/downloads



- No automatic updates (allows version control)
- Downloads an executable installation file
- Administrator access may be required

3) Install as part of Microsoft 365

microsoft.com/en-us/microsoft-365



- Power BI Desktop is included as part of select enterprise Office/Microsoft 365 subscriptions
- If your company uses a compatible version of Microsoft 365, talk to an admin about getting access to Power BI



HEY THIS IS IMPORTANT!

You do **NOT** need to register for a Power BI Pro account to access Power BI Desktop

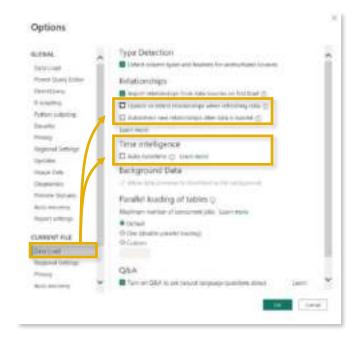
POWER BI SETTINGS

Global > Preview Features



Select **all available preview features** by default (these change with each monthly release)

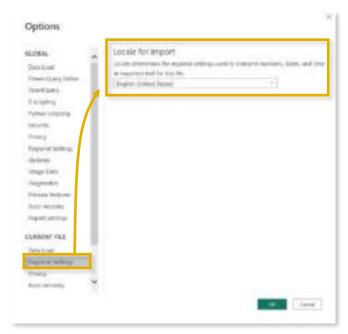
Current File > Data Load



Make sure the following options are **NOT selected**:

- Update or delete relationships when refreshing data
- Autodetect new relationships after data is loaded
- Time Intelligence > Auto date/time

Current File > Regional Settings



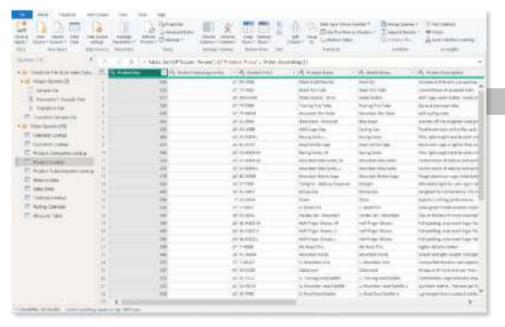
Select "English (United States)" from the dropdown menu (this will align with the data in course project files)



HEY THIS IS IMPORTANT!

Options under CURRENT FILE need to be adjusted every time you open a new Power BI workbook (these settings do not persist across new .pbix files)

Raw data is extracted and transformed in the **Power Query editor**, then loaded to the Power BI "front-end"

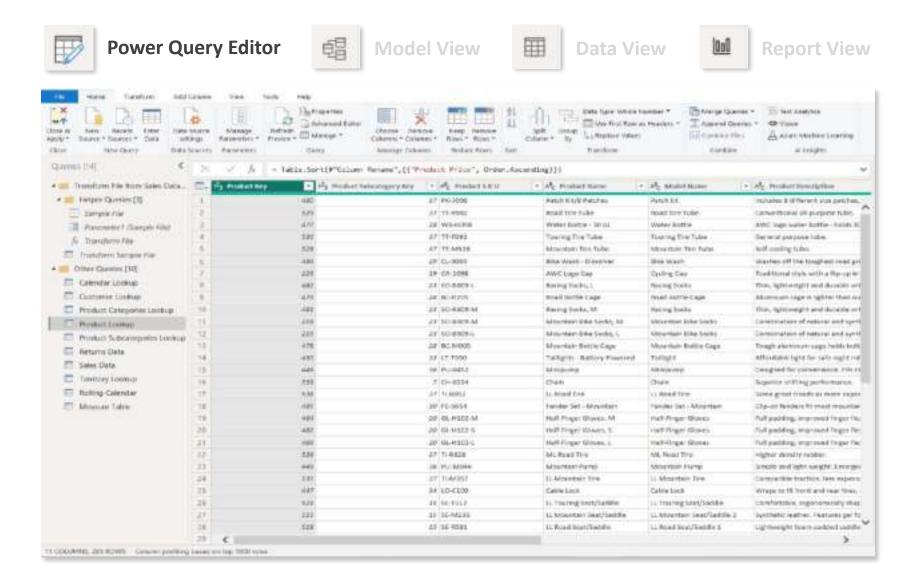




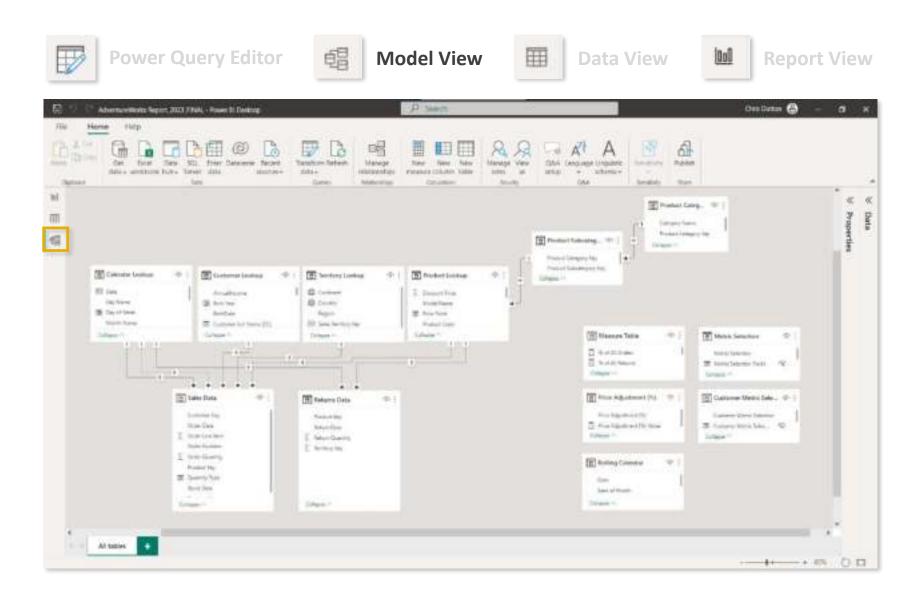
VISUALUTE VINESAUTE Calculations Weekly Streeture Outre by Calegory No. of Persons 619.773 Water Boths: Jillian Make male from her \$10,000 **7869** Mountee Tre Tube \$29,111 146% 13,646 155% Rood Title Take 2173 \$17,285 Special of DE Trafficiant, Navi \$75,499 2392 STURE AME Some Day Sport Elita Meson Blod CHI Personal State - Mosco Rain 507.041 Sport FRE Solenet Mail 1,940 965 JMU Monthly feminer Adventogs (Soders Monthly factors Historian Felde Cause \$140m2 \$1.83M 2,146 Most Criment Frederi Type Most Relational Product Types 166 President Child Colors Tires and Tables Exist Darkfringed Artefact Detail 42 Category Scattor All Decomposition free **Model View Report View Data View**

Oppmus

Power BI "Front-End"



Data is loaded & transformed in the **Power Query Editor**



Data is loaded & transformed in the **Power Query Editor**



Data models are configured in the **Model View**



Power Query Editor



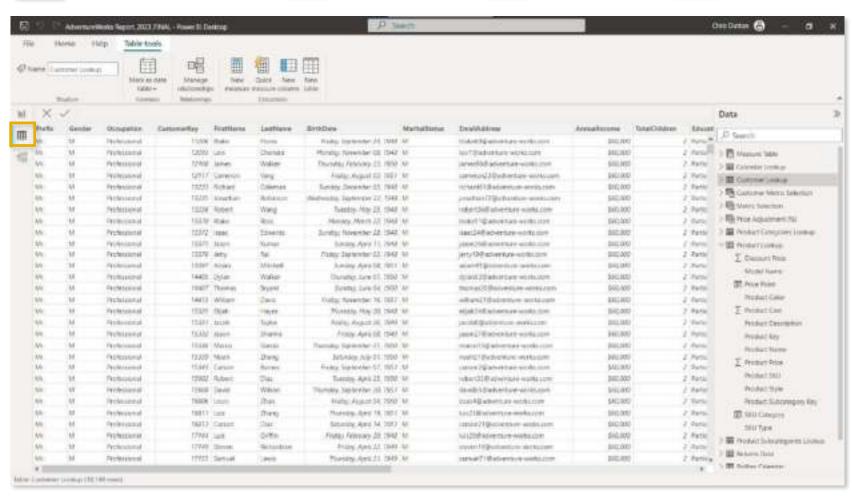
Model View



Data View



Report View



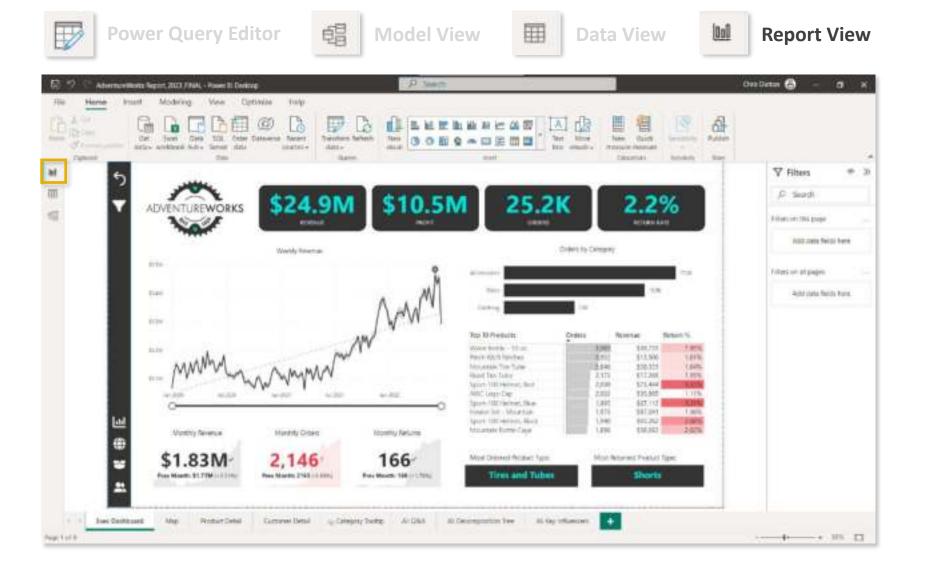
Data is loaded & transformed in the **Power Query Editor**



Data models are configured in the **Model View**



Table features & calculations are added in the **Data View**



Data is loaded & transformed in the **Power Query Editor**



Data models are configured in the **Model View**



Table features & calculations are added in the **Data View**



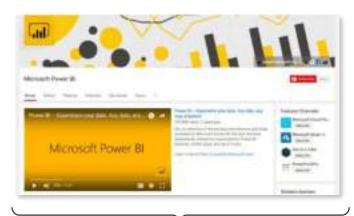
Visuals & reports are designed in the **Report View**

HELPFUL RESOURCES



The **Help** tab includes documentation, training videos, sample files, templates, and links to support blogs and communities





Annual Particles for contrast a fine contrast and a second contras

The Microsoft Power BI blog (powerbi.microsoft.com/blog) publishes monthly summaries to showcase new features

The Microsoft Power BI YouTube Channel publishes demos, feature summaries, and advanced tutorials (check out "Guy in a Cube" too!) **Power BI User Groups** (Power BIUG) are communities of users, which include both local meet-ups and helpful online forums (*pbiusergroup.com*)

MONTHLY UPDATES

Power BI is updated monthly, so you may notice ongoing changes to settings, options, tools, etc. Reference the links below to stay up-to-date on product updates and new feature releases:



Power BI Desktop

https://docs.microsoft.com/en-us/power-bi/fundamentals/desktop-latest-update



Power BI Service

https://docs.microsoft.com/en-us/power-bi/fundamentals/service-whats-new



Power Platform

https://learn.microsoft.com/en-us/dynamics365/release-plans/

CONNECTING & SHAPING DATA

CONNECTING & SHAPING DATA



In this section we'll connect to source files and cover some of the most common techniques for **extracting**, **cleaning**, and **shaping data** to prepare it for modeling and analysis

TOPICS WE'LL COVER:

Intro to Power Query

Data Connectors

The Query Editor

Connection Modes

Data QA & Profiling

Table Transformations

Calendar Tools

Combining Queries

GOALS FOR THIS SECTION:

- Explore Power BI's query editor and understand the role that Power Query plays in the larger BI workflow
- Introduce different types of connectors and connectivity modes available for getting data into Power BI
- Review tools for checking data quality and key profiling metrics like column distribution, empty values, errors and outliers
- Transform tables using text, numerical and date/time tools, pivot and group records, and create new conditional columns
- Practice combining, modifying and refreshing queries

FRONT-END VS. BACK-END



Power BI Desktop essentially has two distinct environments: a front-end and a back-end

- The **front-end** includes the **Data**, **Model** & **Report** views, where most of the modeling, analysis and visualization takes place
- The back-end includes the Power Query Editor, where raw data is extracted, transformed, and loaded to the front-end (ETL)

BACK-END

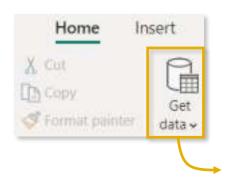
- Connect & extract data using pre-built connectors
- Profile & QA the data to explore, clean and prepare it for modeling and analysis
- Transform & shape tables to add new features, modify values, group records, or sort and filter columns
- Merge or append queries to join and combine them prior to loading to the front-end
- Perform advanced transformations using custom M code (out of scope for this course)

FRONT-END

- Build data models by creating table relationships between primary and foreign keys
- Add calculated measures & columns using Data Analysis Expressions (DAX)
- Design reports to visualize the data and create interactive, dynamic dashboards
- Publish & share your Power BI workbooks using Power BI Service (cloud application)

TYPES OF DATA CONNECTORS





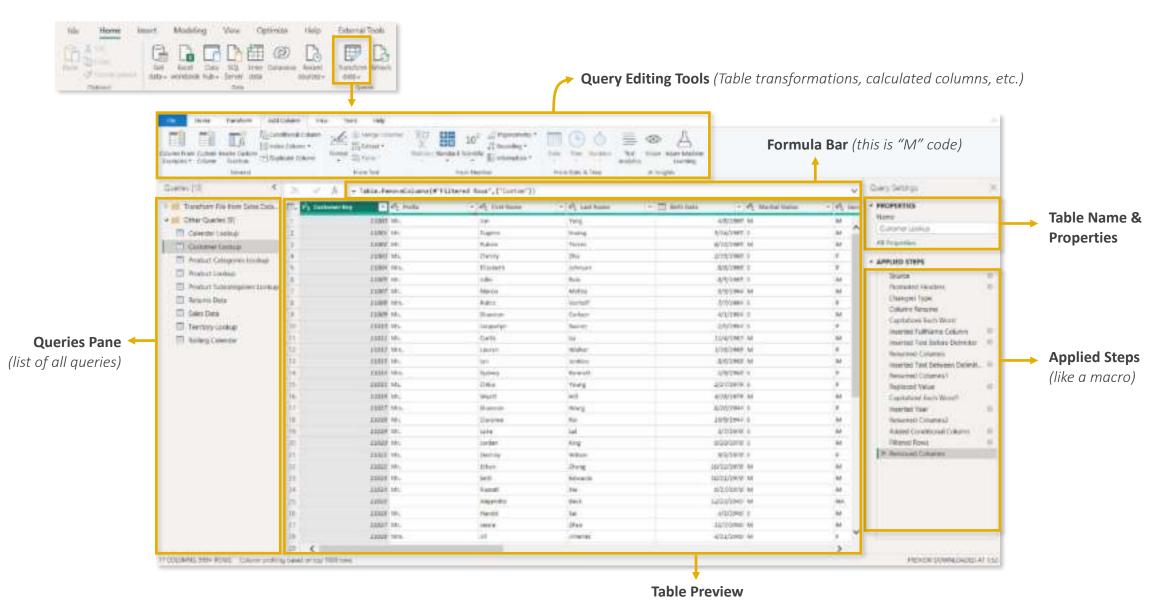


Power BI can connect to virtually **any** type of source data, including (but not limited to):

- Flat files & Folders (csv, text, xlsx, etc.)
- Databases (SQL, Access, Oracle, IBM, etc.)
- **Power Platform** (*Datasets, Datamarts, Dataflows, Dataverse, etc.*)
- Azure (Azure SQL, Analysis Services, Databricks, etc.)
- Online Services (SharePoint, GitHub, Dynamics 365, Google Analytics, Salesforce, Power BI Service, etc.)
- Other (Web feeds, R scripts, Spark, Hadoop, etc.)

POWER QUERY EDITOR

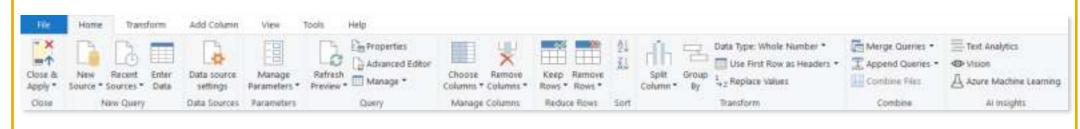




QUERY EDITING TOOLS



The **HOME** tab includes **general settings** and **common table transformation tools**



The ${\it TRANSFORM}$ tab includes tools to ${\it modify}$ ${\it existing}$ ${\it columns}$ (splitting/grouping, transposing, ${\it extracting}$ ${\it text}$, ${\it etc.}$) -

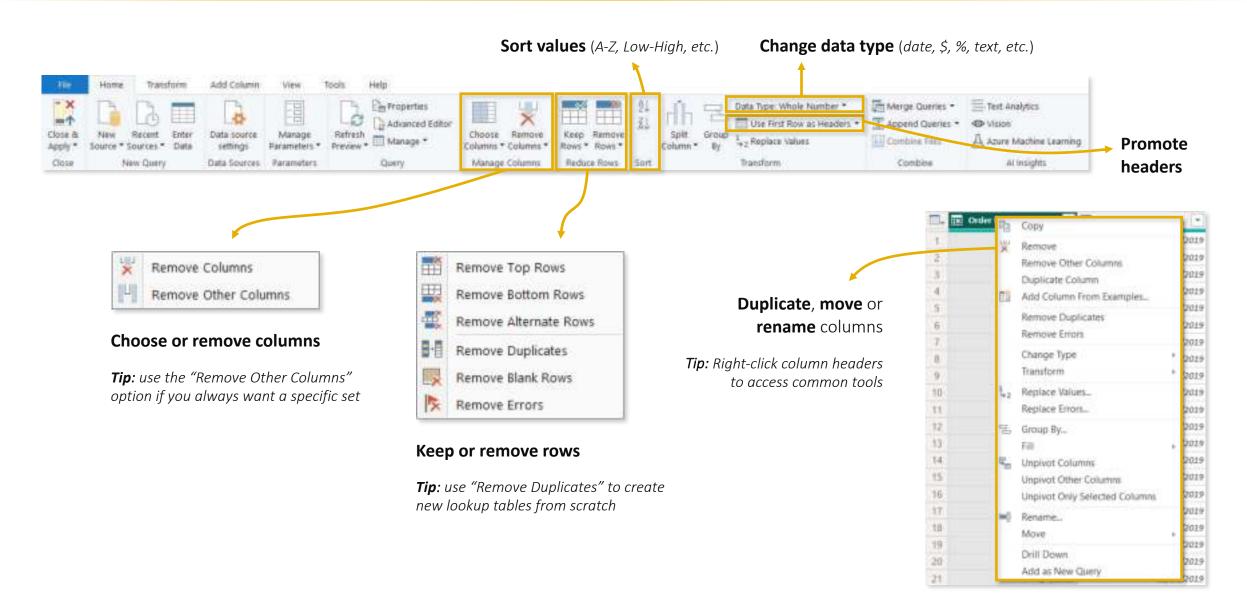


- The **ADD COLUMN** tools **create new columns** (based on conditional rules, text operations, calculations, dates, etc.) -



BASIC TABLE TRANSFORMATIONS





ASSIGNMENT: TABLE TRANSFORMATIONS





From: **Ethan T. Langer** (Analytics Manager)

Subject: Welcome aboard!

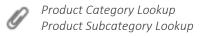
Hello, and welcome to the team!

We're excited that you'll be helping us develop our new internal reports in Power BI. Looks like you've already gotten started, but we have some new data to add to the model.

Could you please create two new queries to connect to the **Product Category Lookup** and **Product Subcategory Lookup** files attached, and help with a few modifications to the product table?

Thanks!

-ETL





Key Objectives

- 1. Create queries to connect to the two new .csv files
- 2. Name your queries **Product Category Lookup** and **Product Subcategory Lookup**
- 3. Confirm that column headers have been promoted and that all data types are correct
- 4. Add a new column to extract all characters before the dash ("-") in the **Product SKU** column, and name it "**SKU Type**"
- 5. Update the **SKU Type** calculation above to return all characters before *second* dash, instead of the first
- 6. Replace zeros (**0**) in the **Product Style** column with "**NA**"
- 7. Close and load to your data model

SOLUTION: TABLE TRANSFORMATIONS





From: **Ethan T. Langer** (Analytics Manager)

Subject: Welcome aboard!

Hello, and welcome to the team!

We're excited that you'll be helping us develop our new internal reports in Power BI. Looks like you've already gotten started, but we have some new data to add to the model.

Could you please create two new queries to connect to the **Product Category Lookup** and **Product Subcategory Lookup** files attached, and help with a few modifications to the product table?

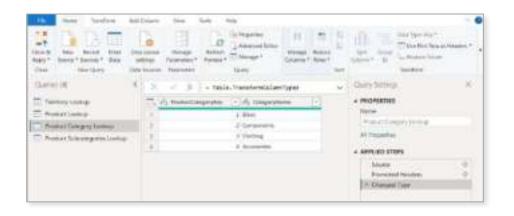
Thanks!

-ETL





Solution Preview





PRO TIP: STORAGE & CONNECTION MODES



Power BI Desktop supports several types of **storage** and **connection modes**:

- Import: Tables are stored in-memory within Power BI and queries are fulfilled by cached data (default)
- **DirectQuery**: Tables are connected directly to the source and queries are executed on-demand at the data source
- Composite Model (Dual): Tables come from a mix of Import and DirectQuery modes, or integrate multiple DirectQuery tables
- Live Connection: Connect to pre-published Power BI datasets in Power BI Service or Azure Analysis Services



Import

- ✓ Dataset is less than 1GB (after compression) & fast performance
- Source data does not change frequently
- ✓ No restrictions on Power Query, data modeling, and DAX functions



DirectQuery

- Dataset is too large to be stored in-memory
- Source data changes frequently and reports must reflect changes
- Company policy states that data can only be accessed from the original source



Composite Model

- ✓ Boost performance by setting appropriate storage for each table
- Combine a DirectQuery model with additional imported data
- ✓ Create a single model from two or more DirectQuery models



Live Connection

- Create one dataset that serves as a central source of truth
- Analyst teams can create different reports from the same source
- Multi-developer teams where one user builds the model and another works on visualization

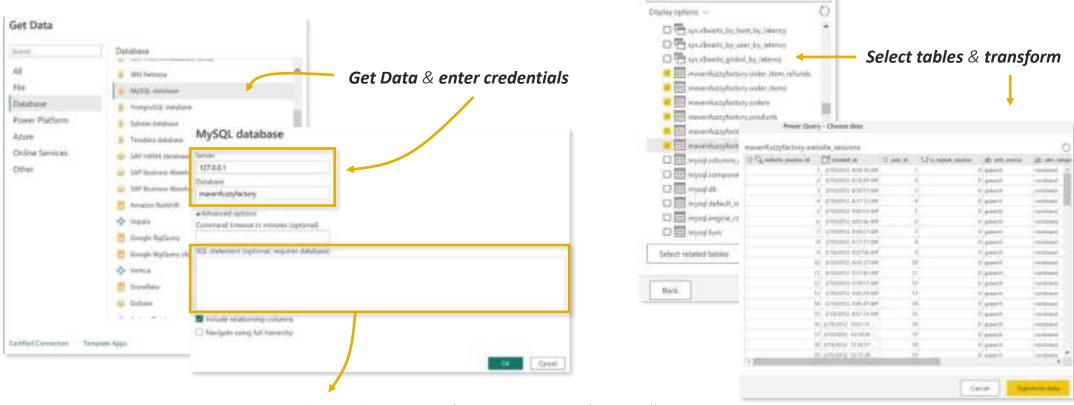
Learn more: https://learn.microsoft.com/en-us/power-bi/connect-data/service-dataset-modes-understand

CONNECTING TO A DATABASE



Power Query can connect to data from various **database sources** including SQL Server, MS Access, MySQL, PostgreSQL, Oracle, SAP, and more

D Greek



Write custom or advanced queries with SQL statements (optional)

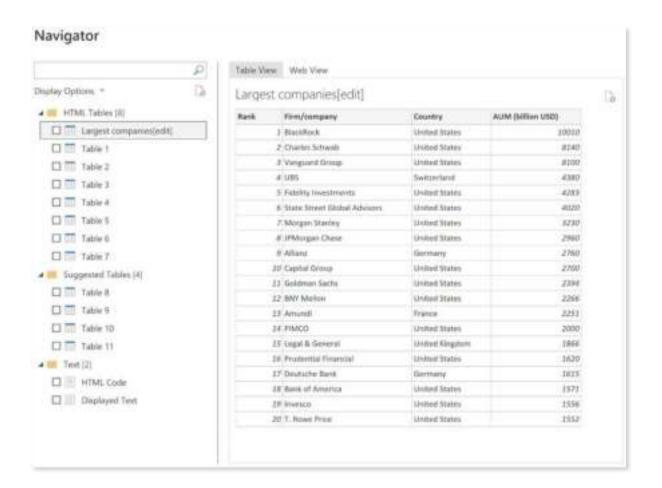
EXTRACTING DATA FROM THE WEB



Power Query includes a native **Web connector** for importing web-hosted files (csv, xlsx, etc.) or scraping URLs for anything that Power Query can identify as a structured table



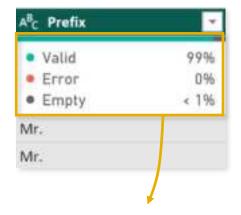
https://en.wikipedia.org/wiki/List_of_asset_management_firms



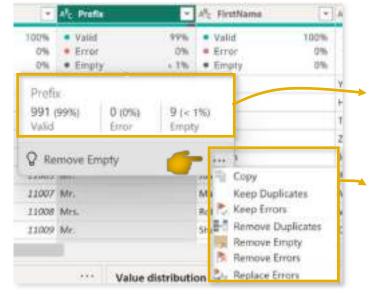
DATA PROFILING: COLUMN QUALITY



Profiling tools like **column quality**, **column distribution**, and **column profile** allow you to explore the quality, composition, and distribution of your data before loading it into the Power BI front-end



Column quality shows the percentage of values within a column that are **valid**, contain **errors**, or are **empty**



Hover over the column quality box to see the **number of records** in each category

Click the **options menu** to remove duplicates, errors or empty values

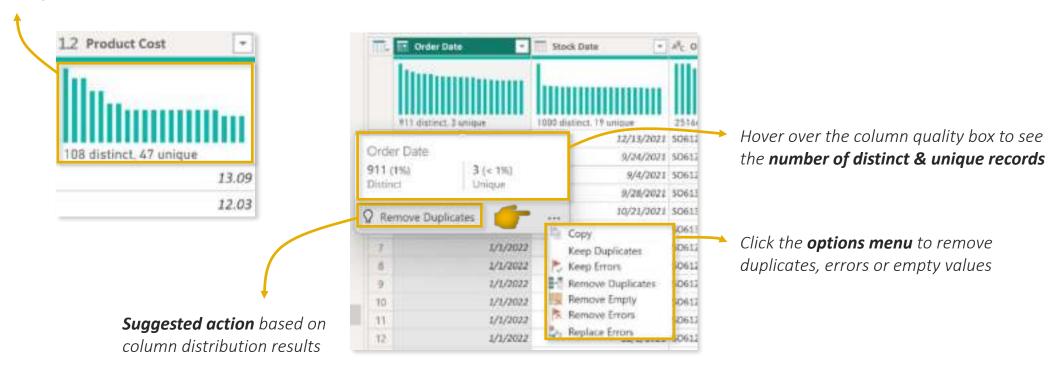


PRO TIP: Profiling tools are a great way to quickly find and address common data quality issues in one place, instead of having to manually apply multiple tools or filters

DATA PROFILING: COLUMN DISTRIBUTION



Column distribution provides a sample distribution of the data in a column



DATA PROFILING: COLUMN PROFILE



Column profile provides a more holistic view of the data in a column, including a sample distribution and profiling statistics

Column statistics provide more detailed profiling metrics, including:

Count = 293

(total number of values in column)

Distinct Count = 119

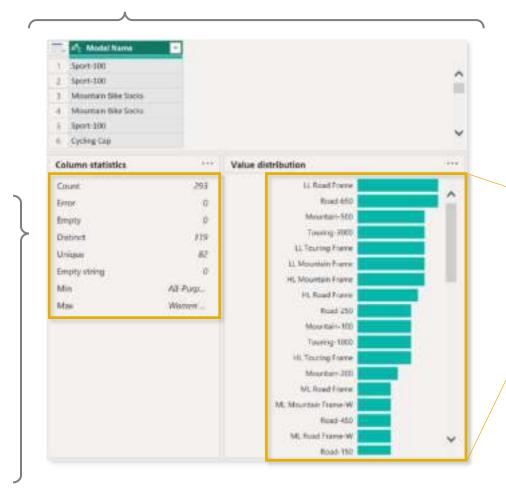
(total number of distinct values, whether they appear once or multiple times)

Unique = 82

(total number of values that appear exactly once)

Min & Max

(lowest and highest observed values) **Note:** Typically only useful for numerical values

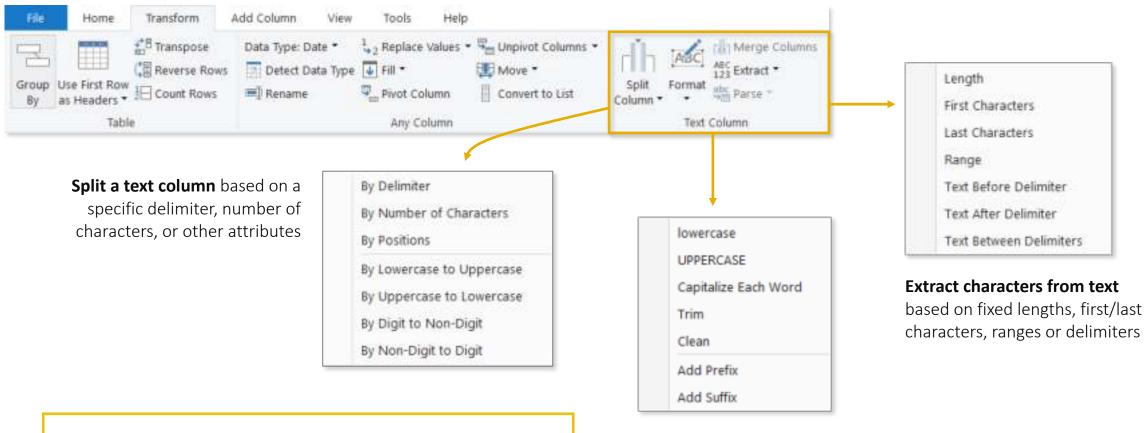




Hover over the value distribution bar for **suggested transformations** and additional options

TEXT TOOLS







HEY THIS IS IMPORTANT!

You can access many tools from both the **Transform** and **Add Column** menus - the difference is whether you want to **ADD** a new column or **OVERWRITE** an existing one

Format a text column to upper, lower or proper case, or add a prefix or suffix

Tip: Use "Trim" to eliminate leading & trailing spaces, or "Clean" to remove non-printable characters

ASSIGNMENT: TEXT TOOLS





From: **Ethan T. Langer** (Analytics Manager)

Subject: **Customer domains**

Hi!

We're looking to better understand where our customers may be coming from, based on their email domains.

Could you please create a new column in the customer table that will allow us do this?

Thanks!

-ETL



Key Objectives

- 1. Duplicate the email address column and name it "**Domain Name**"
- 2. In the new column, remove all text/characters except for the domain name
- 3. Use transformation steps to clean up and capitalize the domain names (i.e. "Adventure Works")
- 4. Save & Apply changes

SOLUTION: TEXT TOOLS





From: **Ethan T. Langer** (Analytics Manager)

Subject: **Customer domains**

Hi!

We're looking to better understand where our customers may be coming from, based on their email domains.

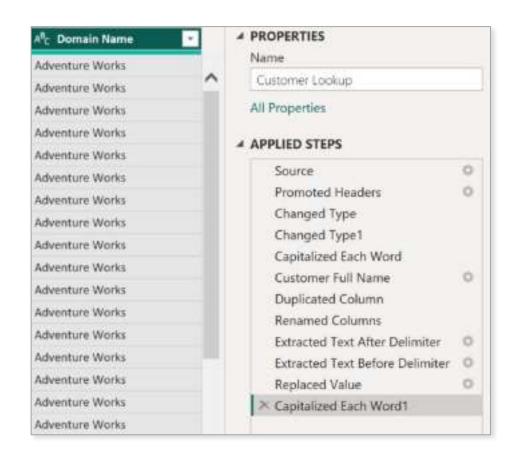
Could you please create a new column in the customer table that will allow us do this?

Thanks!

-ETL

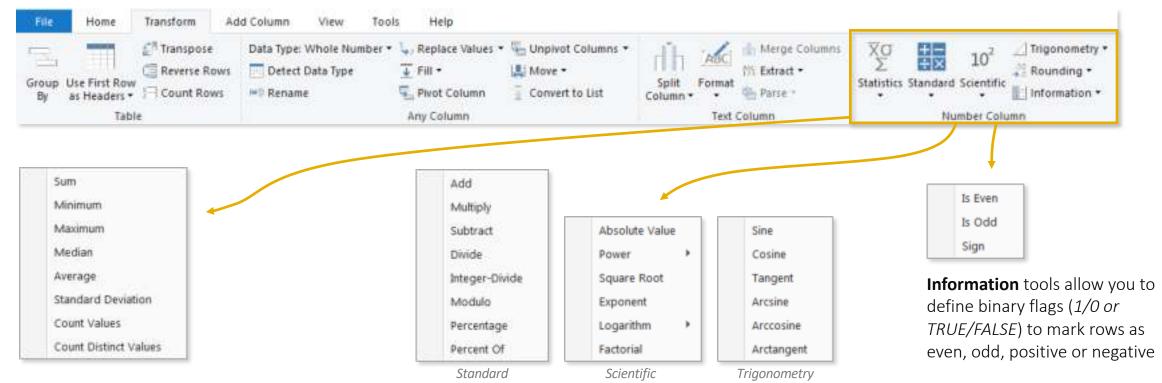


Solution Preview



NUMERICAL TOOLS





Statistics functions allow you to evaluate basic stats for a selected column (sum, min/max, average, count, count distinct, etc.)

Note: These tools return a SINGLE value, and are commonly used to explore a table rather than prepare it for loading

Standard, **Scientific** and **Trigonometry** tools allow you to apply standard operations (addition, multiplication, division, etc.) or more advanced calculations (power, logarithm, sine, tangent, etc.) to each value in a column

Note: Unlike the Statistics tools, these are applied to each row in the table

ASSIGNMENT: NUMERICAL TOOLS





From: **Ethan T. Langer** (Analytics Manager)

Subject: Need some stats for leadership

Hi again,

Leadership is asking us to validate some high-level stats about our products and customers. Can you please help me answer the following questions?

We don't really need to store these values anywhere, so make sure to restore the tables back to their original state once you're done pulling the stats.

Thank you!

-ETL



Key Objectives

- 1. What is our average product cost?
- 2. How many colors do we sell our products in?
- 3. How many distinct customers do we have?
- 4. What is the maximum annual customer income?
- 5. Return the tables to their original state

SOLUTION: NUMERICAL TOOLS





From: **Ethan T. Langer** (Analytics Manager)

Subject: Need some stats for leadership

Hi again,

Leadership is asking us to validate some high-level stats about our products and customers. Can you please help me answer the following questions?

We don't really need to store these values anywhere, so make sure to restore the tables back to their original state once you're done pulling the stats.

Thank you!

-ETL

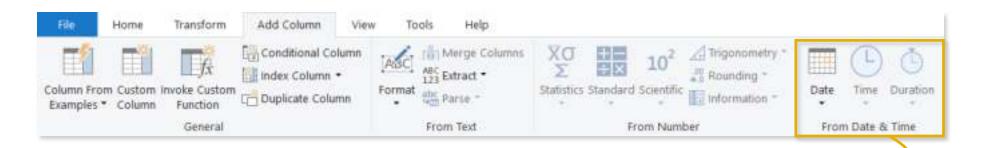


Solution Preview

- 1. What is our average product cost? (\$413.66)
- 2. How many colors do we sell our products in? (10)
- 3. How many distinct customers do we have? (18,148)
- 4. What is the maximum annual customer income? (\$170k)
- 5. Return the tables to their original state

DATE & TIME TOOLS

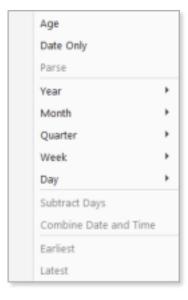




Date & Time tools are relatively straight-forward, and include the following options:

- Age: Difference between the current date and the date in each row
- Date Only: Removes the time component from a date/time field
- **Year/Month/Quarter/Week/Day**: Extracts individual components from a date field (time-specific options include Hour, Minute, Second, etc.)
- **Earliest/Latest:** Evaluates the earliest or latest date from a column as a single value (can only be accessed from the "Transform" menu)

Note: You will almost always want to perform these operations from the "Add Column" menu to build out new fields, rather than transforming an individual date/time column

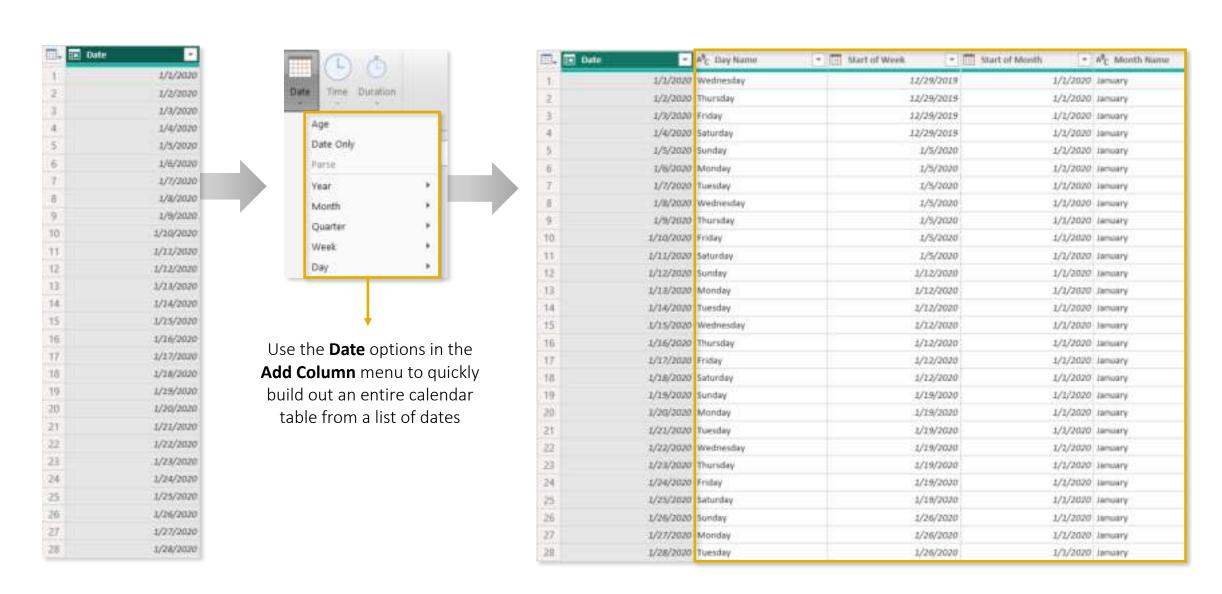




PRO TIP: Load up a table containing a single date column and use Date tools to build out an entire calendar table

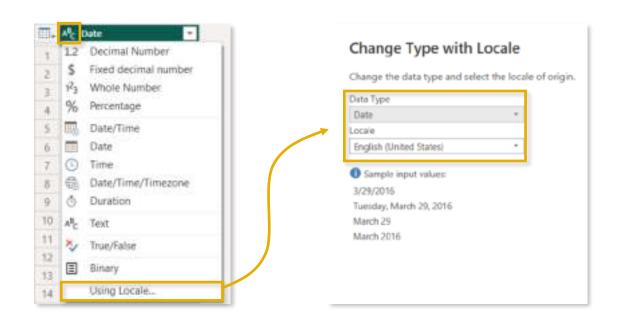
CREATING A CALENDAR TABLE





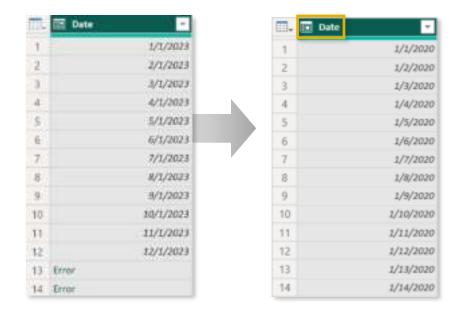
CHANGE TYPE WITH LOCALE





1) Left click the data type icon in the column header and select the Using Locale option

2) Select **Date** as the data type **and English** (**United States**) as the locale for all datasets in this course (regardless of your actual location)

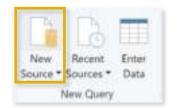


3) Confirm that the data type is correctly recognized. You should see a calendar icon next to the column name in the header and no errors in the column

PRO TIP: ROLLING CALENDARS



Create a new blank query & name it "Rolling Calendar"

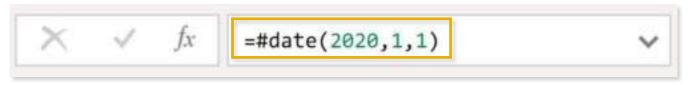


Power Query: New Source > Blank Query



Front end: Get Data > Blank Query

In the formula bar, type a "literal" to generate a start date:



Format as: YYYY, MM, DD

Click the fx icon to **add a custom step**, and enter the following formula to generate a list of dates between the start date and the current day:

```
= List.Dates(
    Source,
    Number.From(DateTime.LocalNow()) - Number.From(Source),
    #duration(1, 0, 0, 0)
)
```

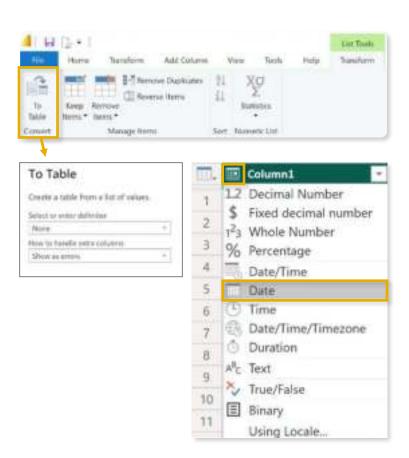
Note: If your first applied step is named something other than "**Source**", use that name in your formula (this is common for non-US users)

PRO TIP: ROLLING CALENDARS



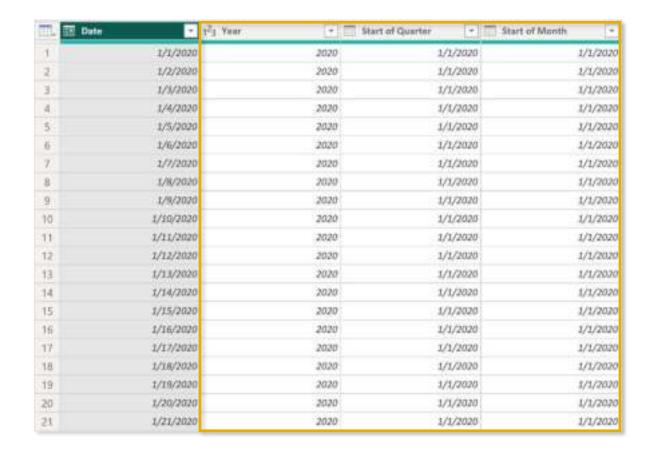
4

Convert the resulting list into a **Table** and set the data type as a **Date**





Rename the column to "**Date**" and add calculated date columns (year, month, quarter, etc.) using the **Add Column** tools



ASSIGNMENT: CALENDAR TABLES





From: **Ethan T. Langer** (Analytics Manager)

Subject: New date fields

Hi,

We need to add a few fields to our calendar table to help us analyze sales trending over time.

Could you please add the following columns when you get a chance?

Thanks!

-ETL



Key Objectives

Add the following columns to the calendar table:

- **1. Month Name** (e.g. "January")
- 2. Month Number (e.g. "1")
- **3.** Start of Year (e.g. "1/1/2020")
- **4. Year** (e.g. "2020")

SOLUTION: CALENDAR TABLES





From: **Ethan T. Langer** (Analytics Manager)

Subject: New date fields

Hi,

We need to add a few fields to our calendar table to help us analyze sales trending over time.

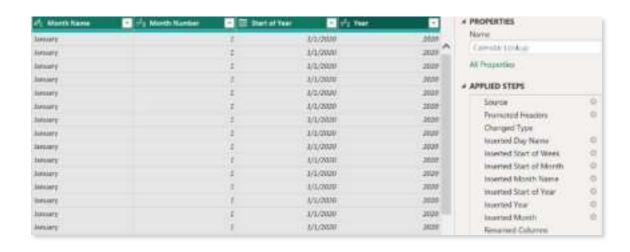
Could you please add the following columns when you get a chance?

Thanks!

-ETL

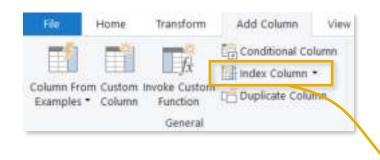


Solution Preview



INDEX COLUMNS





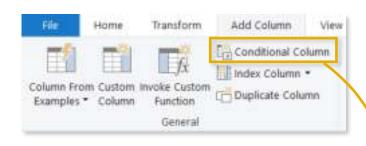
Index Columns contain a list of sequential values that can be used to identify each unique row in a table (*typically starting from 0 or 1*)

These are often used to create **unique IDs** that can be used to form relationships between tables (*more on that later!*)

⊞.	1 ² 2 Index	Order Date -	Stock Date *	A ^B C Order Number -	1 ² 3 Product Key •
1	1	1/1/2020	9/21/2019	5045080	332
2	2	1/1/2020	12/5/2019	5045079	312
3	3	1/1/2020	10/29/2019	5045082	350
4	4	1/1/2020	11/16/2019	5045081	338
5	5	1/2/2020	12/15/2019	5045083	312
6	6	1/2/2020	10/12/2019	5045084	310
7	ÿ	1/2/2020	12/18/2019	5045086	314
0.		1/2/2020	10/9/2019	5045085	312
9	9	1/3/2020	20/3/2019	1045093	312
10	20	1/3/2020	9/29/2019	5045090	320
11	. #1	1/3/2020	12/11/2019	5045088	345
12	12	1/3/2020	10/24/2019	5G4509Z	313
13	11	1/3/2020	12/16/2019	\$045089	.852
14	14	1/3/2020	10/26/2019	5045091	324
15	15	1/3/2020	9/11/2019	5045087	250
16	16	1/3/2020	9/11/2019	5045094	310
17	-17	1/4/2020	10/30/2019	5045096	312
18	-18	1/4/2020	10/30/2019	SG45097	323
19	29	1/4/2020	9/25/2019	5045098	310
20	-20	1/4/2020	12/7/2019	5045095	344

CONDITIONAL COLUMNS

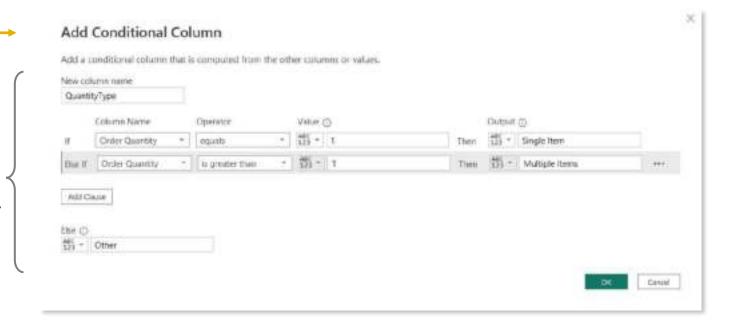




Conditional Columns allow you to define new fields based on logical rules and conditions (IF/THEN statements)

Here we're creating a conditional column named **Quantity Type**, which is based on **Order Quantity**:

- If Order Quantity =1, Quantity Type = "Single Item"
- Else If Order Quantity >1, Quantity Type = "Multiple Items"
- Else; Quantity Type = "Other"



CALCULATED COLUMN BEST PRACTICES



As a best practice, table transformations and column calculations should ideally happen as close to the original data source as possible, to optimize performance and speed









Published Reports

UPSTREAM

DOWNSTREAM

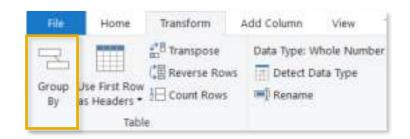


HEY THIS IS IMPORTANT!

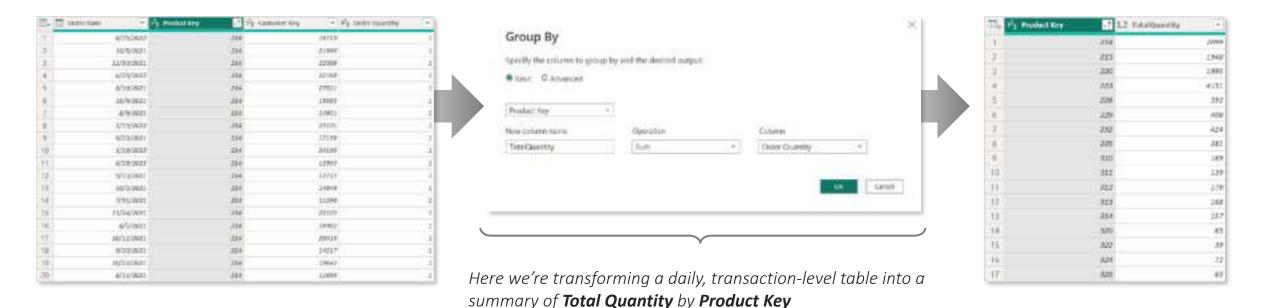
This is not a strict rule or requirement but can significantly impact performance for very large or complex data models. Where you define calculations often depends on several factors (accessibility, complexity, business requirements, etc.), so we will practice creating columns using both Power Query and the Power BI front-end (DAX) throughout this course

GROUPING & AGGREGATING





Group By allows you to aggregate data at a different level or "grain" (i.e. group daily records into monthly, aggregate transactions by store, etc.)



NOTE: Any fields not specified in the Group By settings are lost

GROUPING & AGGREGATING





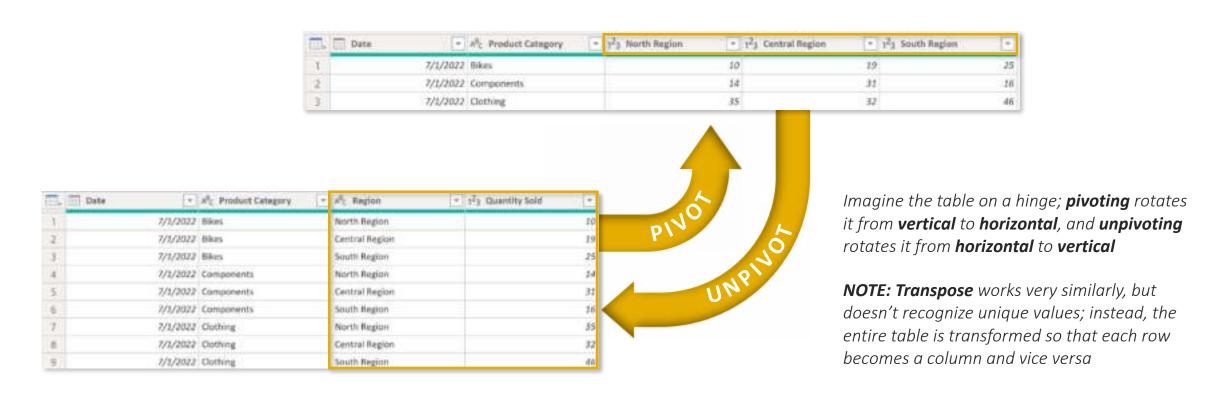
This time we're transforming the daily, transaction-level table into a summary of **Total Quantity** grouped by both **Product Key** and **Customer Key** (using the "Advanced" option)

NOTE: This is like creating a PivotTable in Excel and pulling in **Sum of Order Quantity** with **Product Key** and **Customer Key** as row labels

PIVOTING & UNPIVOTING

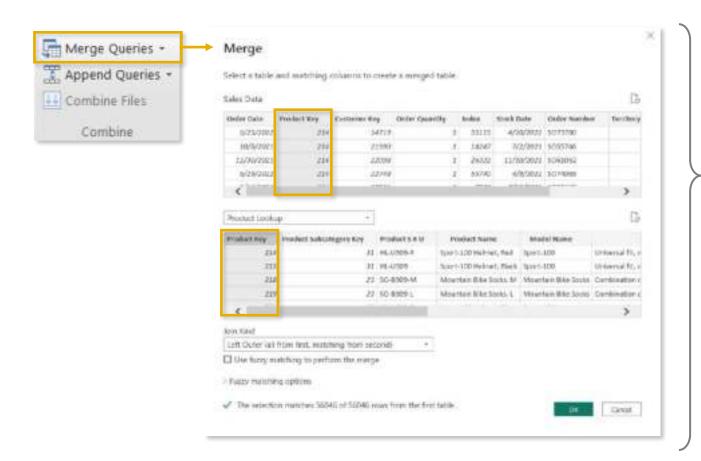


Pivoting describes the process of turning **distinct row values into columns**, and **unpivoting** describes the process of turning **distinct columns into rows**



MERGING QUERIES





Merging queries allows you to **join tables** based on a common column (like a lookup in Excel)

In this case we're merging the **Sales Data** table with the **Product Lookup** table, which share a common **Product Key** column

NOTE: Merging **adds columns** to an existing table/query



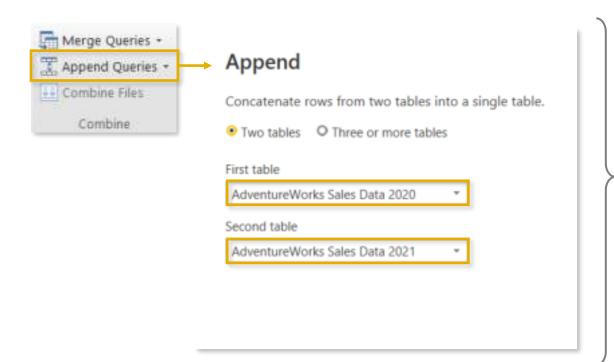
HEY THIS IS IMPORTANT!

Just because you can merge tables, doesn't mean you should!

In many cases, it's better to keep tables separate and define **relationships** between them in the data model (*more on that soon!*)

APPENDING QUERIES





Appending queries allows you to **combine** or **stack** tables sharing the exact same column structure and data types

Here we're appending the **AdventureWorks Sales 2020** table to the **AdventureWorks Sales 2021** table, which is valid since they share identical table structures

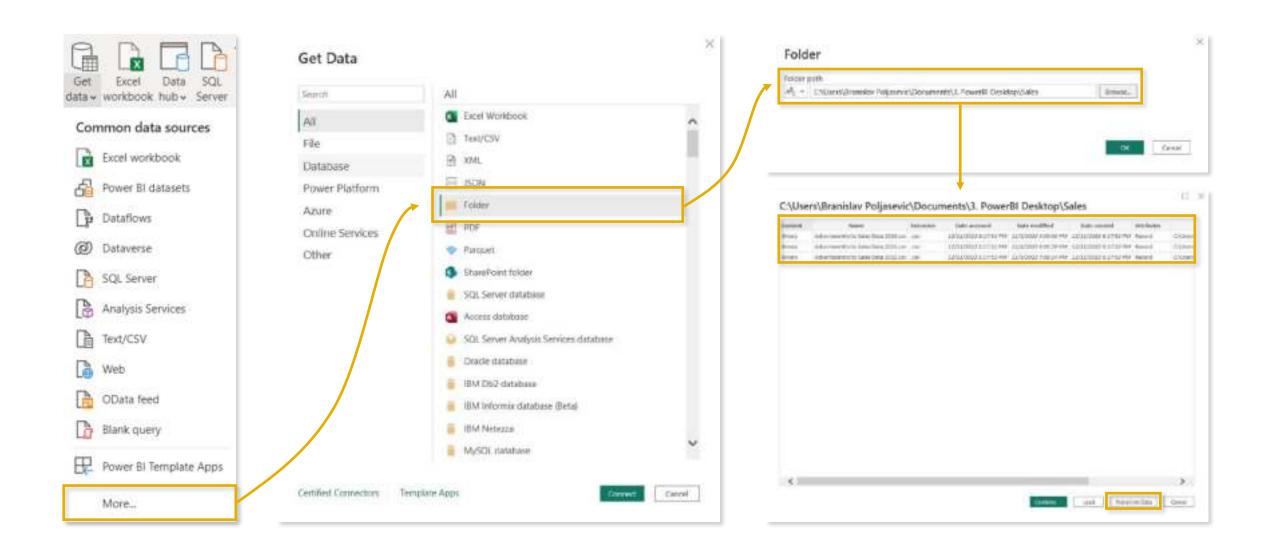
NOTE: Appending **adds rows** to an existing table/query



PRO TIP: Use the **Folder** option (Get Data > More > Folder) to **append all files within a specified folder** (assuming they share the same structure); as you add new files, simply refresh the query and they will automatically append!

PRO TIP: APPENDING FILES FROM A FOLDER

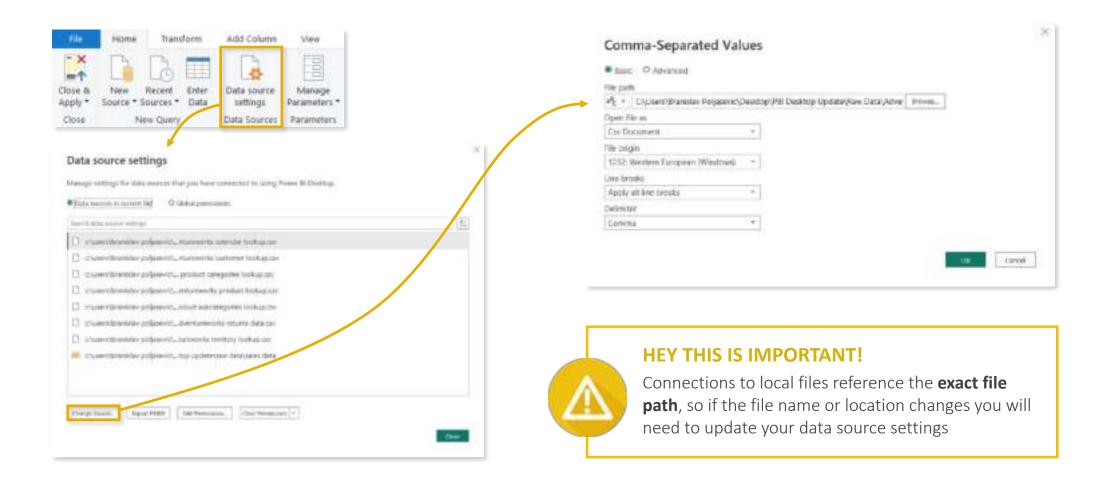




DATA SOURCE SETTINGS



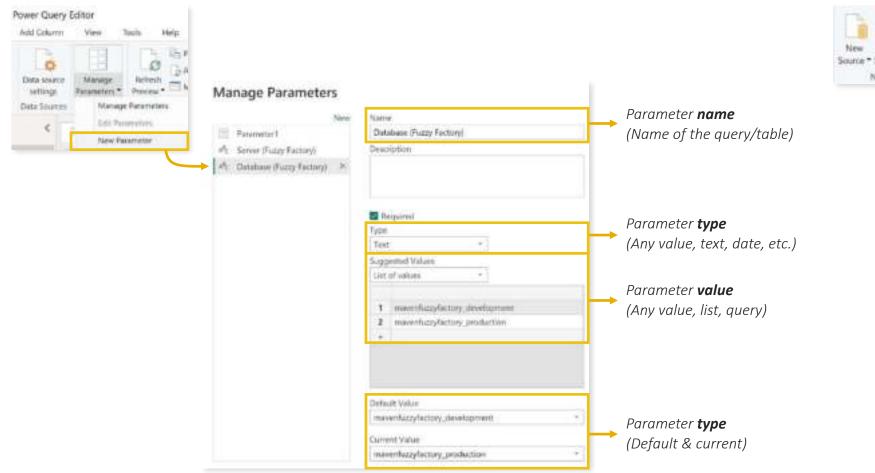
Data Source Settings allow you to manage existing data connections, file paths and permissions

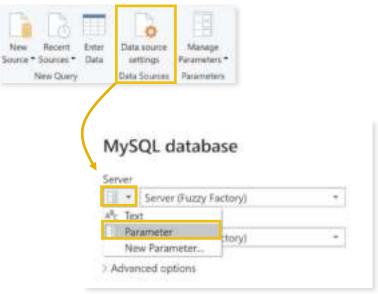


PRO TIP: DATA SOURCE PARAMETERS



Use **parameters** to dynamically manage and update connection paths in the Power Query editor

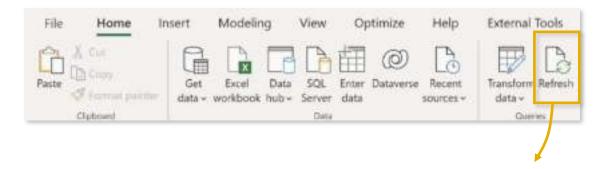




Update Server & Database connection text values **with parameters**

REFRESHING QUERIES



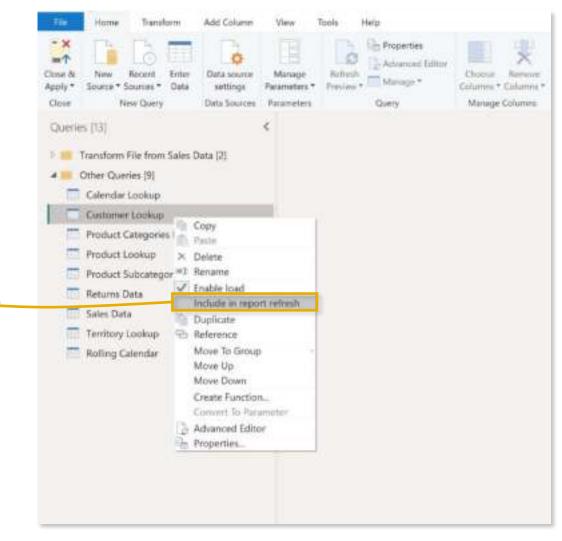


By default, *all queries* will refresh when you use the **Refresh** command from the **Home** tab

From the Query Editor, uncheck **Include in report refresh** to exclude individual queries from the refresh

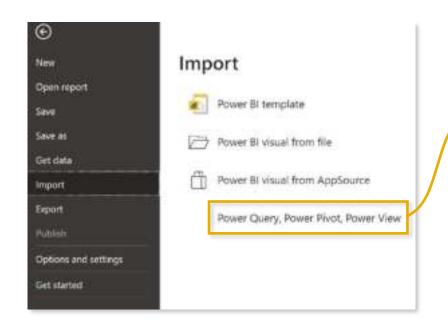


PRO TIP: Exclude queries from refresh that don't change often (like lookups or static data tables)



PRO TIP: IMPORTING EXCEL MODELS





Already have a fully-built model in Excel?

You can import models built in Excel directly into Power BI Desktop using: *Import > Power Query, Power Pivot, Power View*

Imported models retain the following:

- Data source connections and queries
- Query editing procedures and applied steps
- Table relationships, hierarchies, field settings, etc.
- All calculated columns and DAX measures



PRO TIP: If you are more comfortable working in Excel, build your models there first then import to Power BI!

POWER QUERY BEST PRACTICES





Get organized before connecting and loading data

• Define clear and intuitive table/query names from the start, and establish an organized file/folder structure if you are working with local flat files to avoid changes to file names or paths



Disable report refresh for any static data sources

• There's no need to constantly refresh data sources that don't change, like lookups or static data tables



When working with large tables, only load the data you need

 Don't include hourly data when you only need daily, or transaction-level data when only need a product-level summary (extra data will only slow your report down!)



CREATING A DATA MODEL



In this section we'll cover **foundational data modeling topics** like normalization, fact and dimension tables, primary and foreign keys, relationship cardinality and filter flow

TOPICS WE'LL COVER:

Data Modeling 101

Normalization

Facts & Dimensions

Primary & Foreign Keys

Cardinality

Filter Flow

Common Schemas

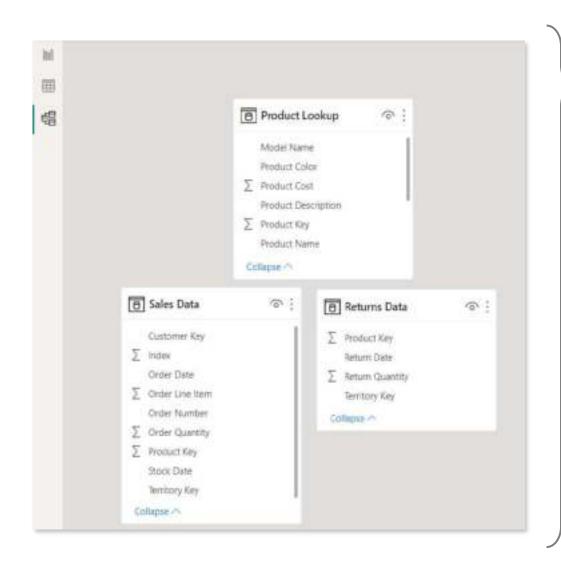
Hierarchies

GOALS FOR THIS SECTION:

- Understand the basic principles of data modeling, including normalization, fact & dimension tables and common schemas
- Create table relationships using primary and foreign keys, and discuss different types of relationship cardinality
- Configure report filters and trace filter context as it flows between related tables in the model
- Explore data modeling options like hierarchies, data categories and hidden fields

WHAT IS A DATA MODEL?





This **IS NOT** a data model

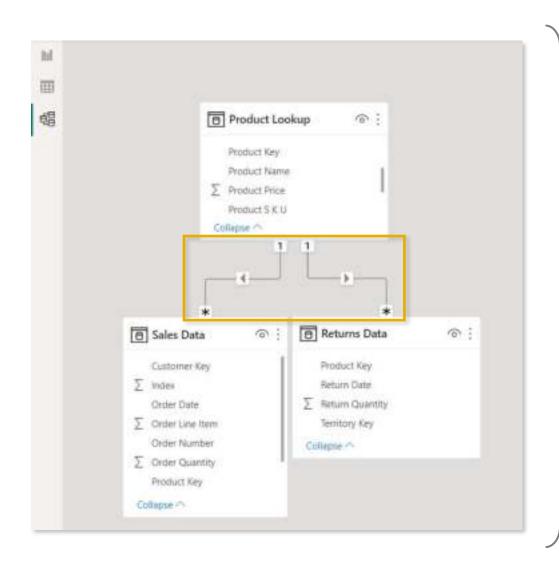


- This is a collection of independent tables, which share no connections or relationships
- If you tried to visualize Orders and Returns by Product, this is what you'd get

ProductName	OrderQuantity	ReturnQuantity
All-Purpose Bike Stand	84,174	1,828
AWC Logo Cap	84,174	1,828
Bike Wash - Dissolver	84,174	1,828
Cable Lock	84,174	1,828
Chain	84,174	1,828
Classic Vest, L	84,174	1,828
Classic Vest, M	84,174	1,828
Classic Vest, S	84,174	1,828
Fender Set - Mountain	84,174	1,828
Total	84,174	1,828

WHAT IS A DATA MODEL?





This **IS** a data model!



- The tables are connected via relationships, based on a common field (Product Key)
- Now Sales and Returns data can be filtered using fields from the Product Lookup table!

ProductName	OrderQuantity	ReturnQuantity
All-Purpose Bike Stand	234	8
AWC Lago Cap	4,151	46
Bike Wash - Dissoiver	1,706	25
Classic Vest, L	182	4
Classic Vest, M	182	7
Classic Vest, S	157	5
Fender Set - Mountain	3,960	54
Half-Finger Gloves, L	840	16
Half-Finger Gloves, M	918	16
Total	84,174	1,828

DATABASE NORMALIZATION



Normalization is the process of organizing the tables and columns in a relational database to reduce redundancy and preserve data integrity. It's commonly used to:

- Eliminate redundant data to decrease table sizes and improve processing speed & efficiency
- Minimize errors and anomalies from data modifications (inserting, updating or deleting records)
- Simplify queries and structure the database for meaningful analysis



In a normalized database, each table should serve a **distinct** and **specific** purpose (i.e. product information, transaction records, customer attributes, store details, etc.)

date *	product_id =	quantity -	product_brand +	product_name -	product sku +	product_weight -
1/1/1997	869	5	Nationeel	Nationeel Grape Fruit Roll	52182117179	17
1/7/1997	869	2	Nationeel	Nationeel Grape Fruit Roll	52382137179	17
1/3/1997	1.	4	Washington	Washington Berry Juice	90748583674	8.39
1/1/1997	1472	3	Fort West	Fort West Fudge Cookies	37276054024	8.28
1/6/1997	1472	2	Fort West	Fort West Fudge Cookies	37276054024	8.28
1/5/1997	2	-4	Washington	Washington Mango Drink	96516502499	7.42
1/1/1997	76	4	Red Spade	Red Spade Sliced Chicken	62054644227	18.1
1/1/1997	76	2	Red Spade	Red Spade Sliced Chicken	62054644227	18.1
1/5/1997	3	2	Washington	Washington Strawberry Drink	58427771925	13.1
1/7/1997	3	2	Washington	Washington Strawberry Drink	58427771925	13.1
1/1/1997	320	3	Excellent	Excellent Cranberry Juice	36570182442	16.4

Models that aren't normalized contain **redundant**, **duplicate data**. In this case, all of the product-specific fields could be stored in a separate table containing a unique record for each **product id**

This may not seem critical now, but minor inefficiencies can become major problems at scale!

FACT & DIMENSION TABLES



Data models generally contain two types of tables: **fact** ("data") tables, and **dimension** ("lookup") tables:

- Fact tables contain numerical values or metrics used for summarization (sales, orders, transactions, pageviews, etc.)
- Dimension tables contain descriptive attributes used for filtering or grouping (products, customers, dates, stores, etc.)

date -	product_id *	quantity -
1/1/1997	869	5
1/1/1997	1472	3
1/1/1997	76	4
1/1/1997	320	3
1/1/1997	4	4
1/1/1997	952	4
1/1/1997	1222	4
1/1/1997	517	4
1/1/1997	1359	4
1/1/1997	357	4
1/1/1997	1426	5
1/1/1997	190	4
1/1/1997	367	4
1/1/1997	250	5
1/1/1997	600	4
1/1/1997	702	5

This **Fact** table contains **quantity** values, along with **date** and **product_id** fields

date -	day_of_month -	month -	year *	weekday -	week_of_year -	week_ending •	month_name =	quarter =
1/1/1997	1	1	1997	Wednesday	1	1/5/1997	January	Q1
1/2/1997	2	1	1997	Thursday	1	1/5/1997	January	Q1
1/3/1997	3	1	1997	Friday	1	1/5/1997	January	Q1
1/4/1997	4	1	1997	Saturday	1	1/5/1997	January	Q1
1/5/1997	5	1	1997	Sunday	2	1/5/1997	January	Q1
1/6/1997	6	1	1997	Monday	2	1/12/1997	January	Q1

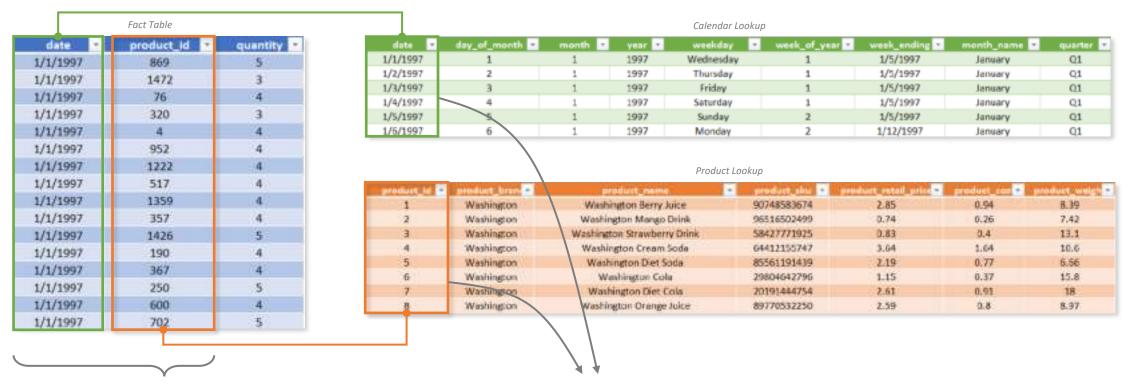
This **Calendar Lookup** table contains attributes about each **date** (month, year, quarter, etc.)

product ld .	preduct brance	product_name	1-1	product shu in	product retail price -	product_res =	product_weigh
1	Washington	Washington Berry Juice		90748583674	2.85	0.94	6.39
2	Washington	Washington Mango Drink		95516502499	0.74	0.26	7.42
3	Washington	Washington Strawberry Drink		58427771925	0.83	0.4	13.1
4	Washington	Washington Cream Sode		64412155747	3.64	1.64	10.6
5	Washington	Washington Diet Soda		85561191439	2.19	0.77	6.66
6	Washington	Washington Cola		29804642796	1.15	0.37	15.8
7	Washington	Washington Diet Cols		20191444754	2.61	0.91	18
8	Washington	Washington Orange Juice		89770532250	2.59	0.8	8.97

This **Product Lookup** table contains attributes about each **product id** (brand, SKU, price, etc.)

PRIMARY & FOREIGN KEYS





These are **foreign keys** (FK)

They contain multiple instances of each value, and relate to **primary keys** in dimension tables

These are **primary keys** (PK)

They uniquely identify each row of the table, and relate to **foreign keys** in fact tables

RELATIONSHIPS VS. MERGED TABLES





Can't I just merge queries or use lookup functions to **pull everything into one single table**?

- Anonymous confused man

Original **Fact Table** fields

Attributes from **Calendar Lookup** table

Attributes from **Product Lookup** table

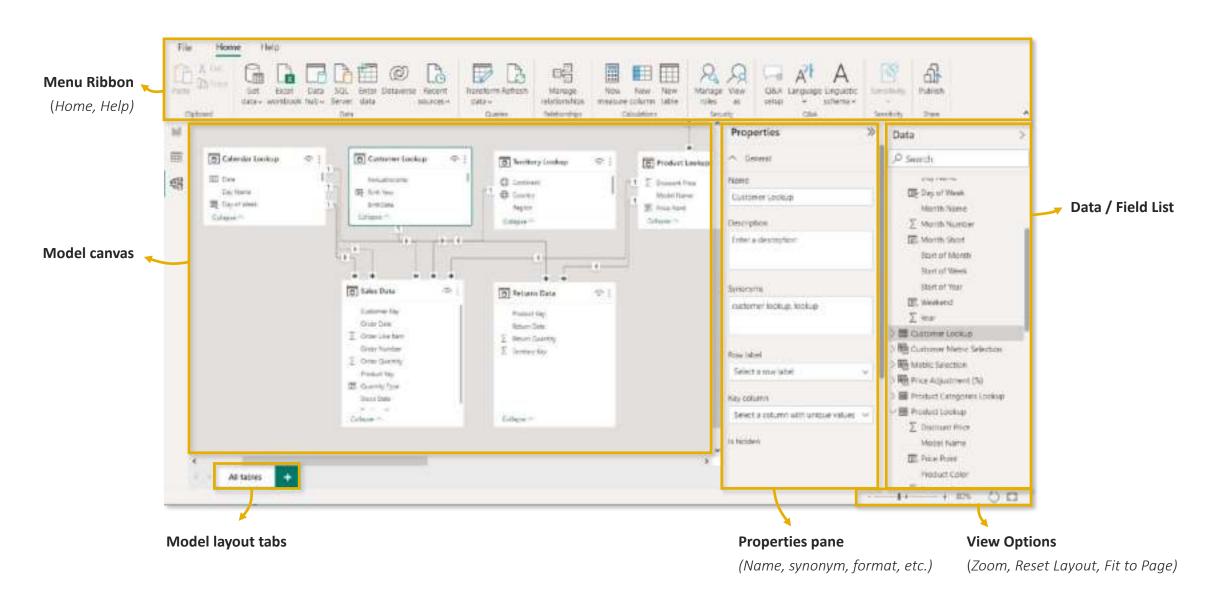
date *	product_id -	quantity -	day of month	month *	year -	weekday -	month_name -	quarter *	product_brand *	product_name	product_sku *	product_weight *
1/1/1997	860	5	A CONTRACTOR OF THE PARTY OF TH	1	1997	Wednesday	January	Q1	Nationsel	Nationeel Grape Fruit Roll	52382137170	17
1/7/1997	869	2	7	1	1997	Tuesday	January	Q1	Nationeel	Nationeel Grape Fruit Roll	52382137179	17
1/3/1997	1	4	3	1	1997	Friday	January	QI	Washington	Washington Berry Juice	90748583674	8.19
1/1/1997	1472	3	1	1	1997	Wednesday	January	Q1	Fort West	Fort West Fudge Cookies	37276054024	8,28
1/6/1997	1472	2	6	1	1997	Monday	January	Q1	Fort West	Fort West Fudge Cookies	37276054024	8.28
1/5/1997	2	-4	5	3	1997	Sunday	January	01	Washington	Washington Mango Drink	96516502499	7.42
1/1/1997	76	4	1	1	1997	Wednesday	January	Q1	Red Spade	Red Spade Sliced Chicken	62054644227	18.1
1/1/1997	76	2	1	1	1997	Wednesday	January	Q1	fied Spade	Red Spade Sliced Chicken	62054644227	18.1
1/5/1997	3	2	5	1	1997	Sunday	January	Q1	Washington	Washington Strawberry Drink	58427771925	13.1
1/7/1997	3	2	7	1	1997	Tuesday	January	Q1	Washington	Washington Strawberry Drink	58427771925	13.1
1/1/1997	320	3	1	1	1997	Wednesday	January	Q1	Excellent	Excellent Cranberry Juice	36570182442	16.4

You can, but it's extremely inefficient!

Merging tables creates **redundancy** and often requires **significantly more memory and processing power** to analyze compared to a relational model with multiple small tables

THE MODEL VIEW

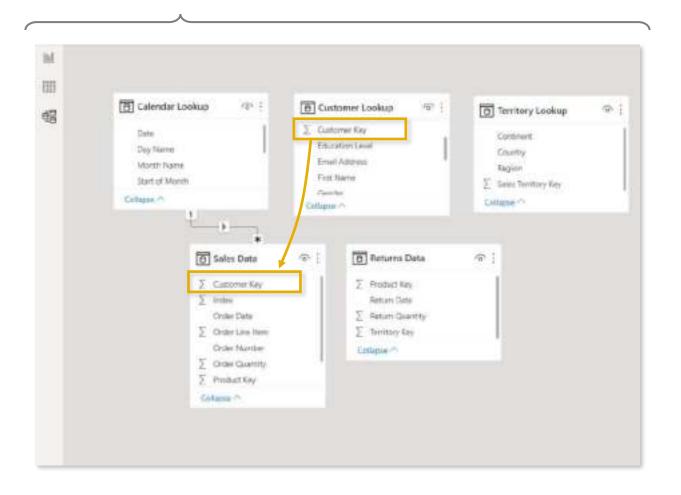




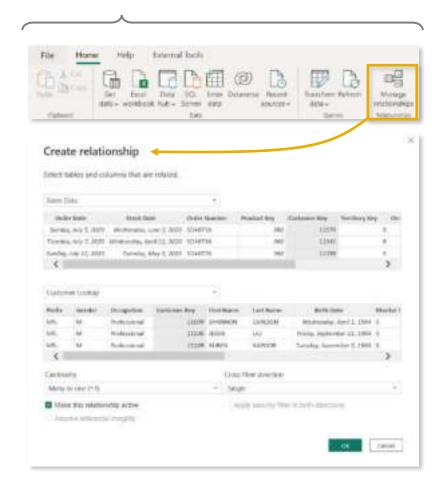
CREATING TABLE RELATIONSHIPS



OPTION 1: Click and drag to connect primary and foreign keys within the **Model** view

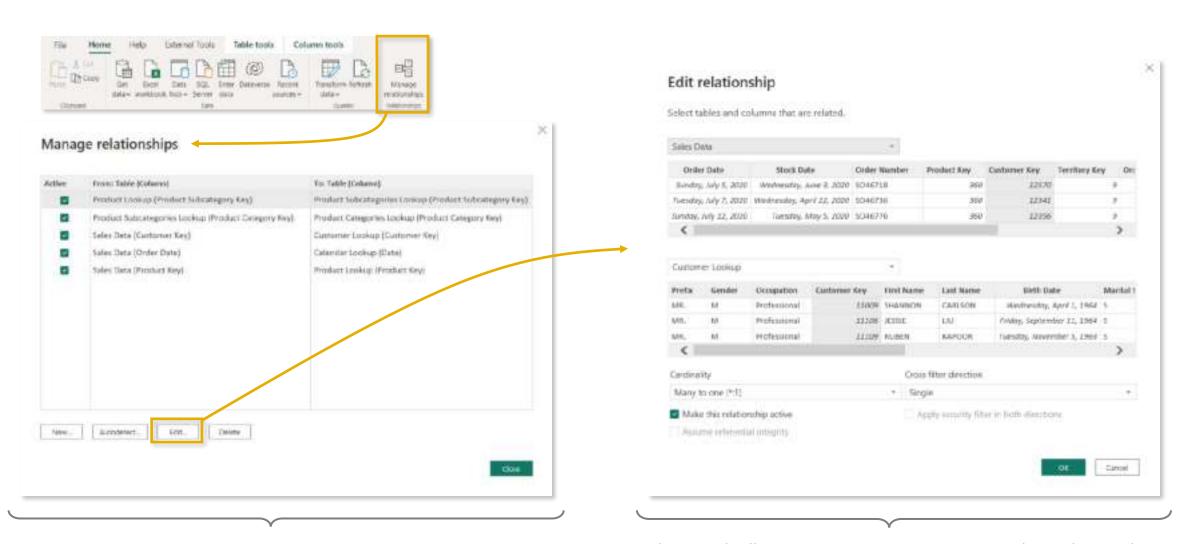


OPTION 2: Add or detect relationships using the **Manage Relationships** dialog box



MANAGING & EDITING RELATIONSHIPS



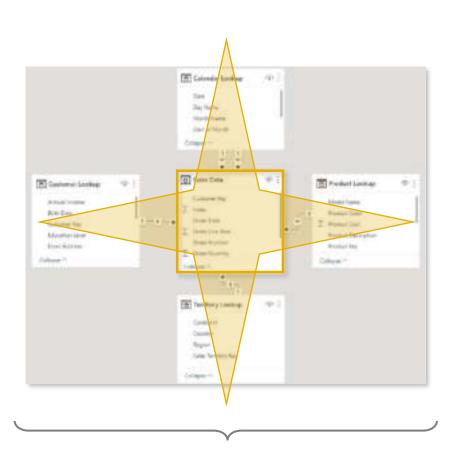


Launch the **Manage Relationships** dialog box or double-click a relationship to modify it

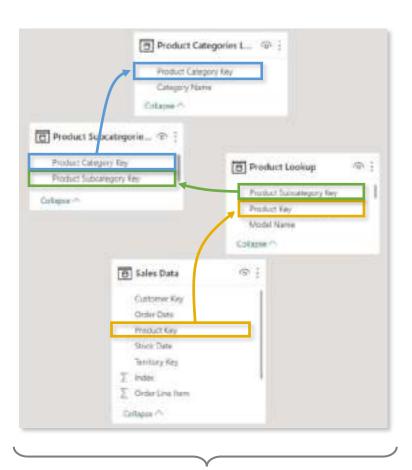
Editing tools allow you to **activate or deactivate** relationships and manage **cardinality** and **filter direction** – more on that soon!

STAR & SNOWFLAKE SCHEMAS





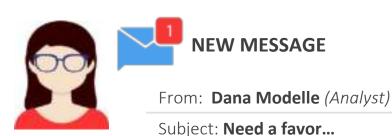
A **star schema** is the simplest and most common type of data model, characterized by a single fact table surrounded by related dimension tables



A **snowflake schema** is an extension of a star, and includes relationships between dimension tables and related sub-dimension tables

ASSIGNMENT: TABLE RELATIONSHIPS





Hey there,

Ethan shared the data model you've been working on, and we might have an issue...

Last night I left my laptop open, and my cat Dennis somehow got his paws on our model. Now all the relationships are gone!

Could you please rebuild the model, including all three product tables? I owe you one!

-Dana



Key Objectives

- 1. Delete all existing table relationships
- Create a star schema by creating relationships between the Sales, Calendar, Customer, Product and Territories tables
- Connect all three product tables (Product, Subcategory, Category) in a snowflake schema
- 4. Use the matrix visual to confirm that you can filter Order Quantity values using fields from each dimension table

SOLUTION: TABLE RELATIONSHIPS





From: **Dana Modelle** (Analyst)

Subject: **Need a favor...**

Hey there,

Ethan shared the data model you've been working on, and we might have an issue...

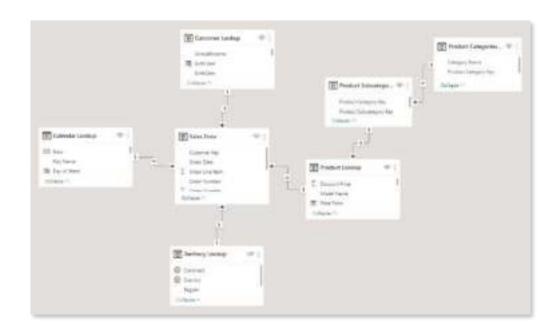
Last night I left my laptop open, and my cat Dennis somehow got his paws on our model. Now all the relationships are gone!

Could you please rebuild the model, including all three product tables? I owe you one!

-Dana

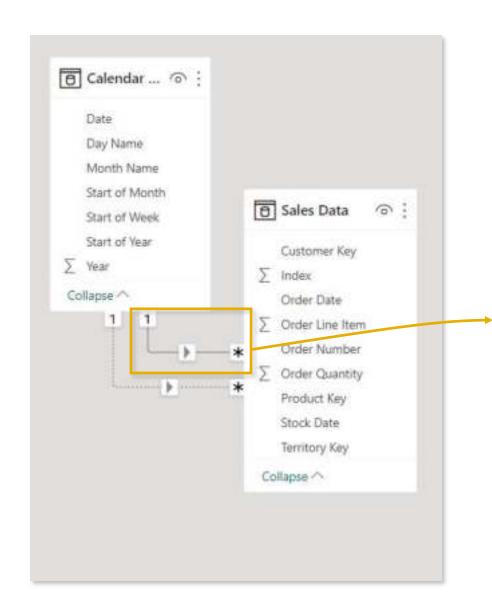


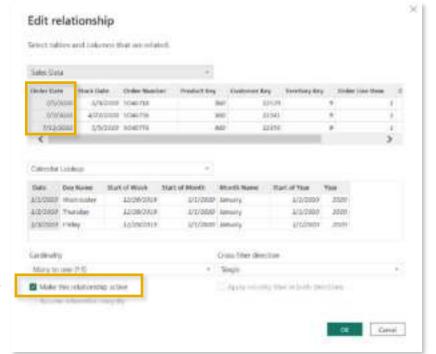
Solution Preview

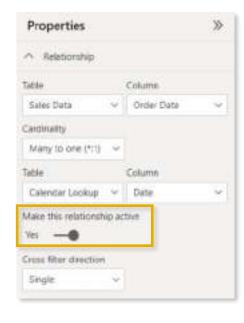


PRO TIP: ACTIVE & INACTIVE RELATIONSHIPS







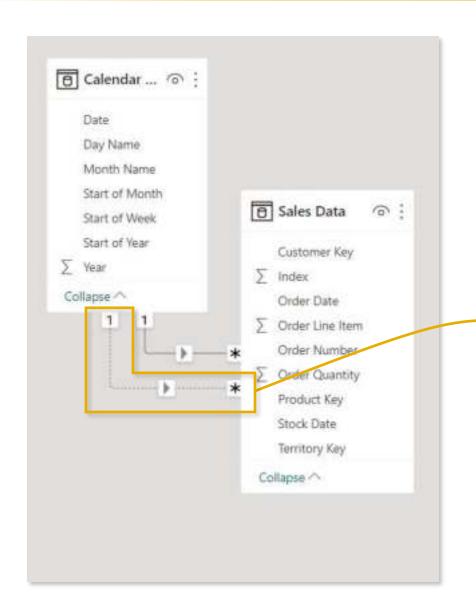


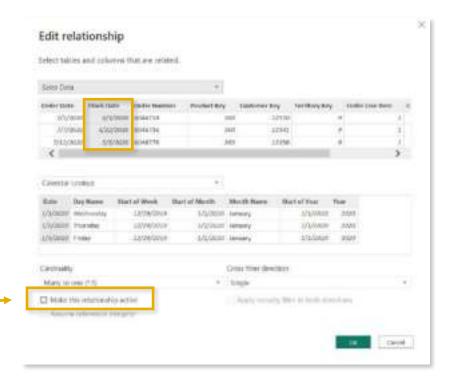
The **Sales Data** table contains two date fields (**Order Date** & **Stock Date**), but there can only be **one active relationship** to the Date key in the Calendar table

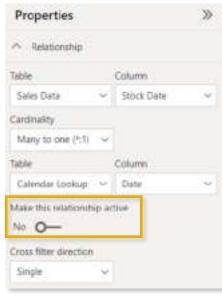
You can set relationships to active or inactive from either the **Edit Relationships** dialog box or the **Properties** (you must deactivate one before activating another)

PRO TIP: ACTIVE & INACTIVE RELATIONSHIPS







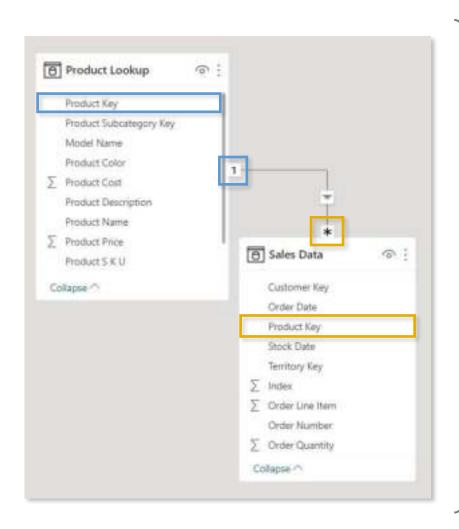


The **Sales Data** table contains two date fields (**Order Date** & **Stock Date**), but there can only be **one active relationship** to the Date key in the Calendar table

You can set relationships to active or inactive from either the **Edit Relationships** dialog box or the **Properties** (you must deactivate one before activating another)

RELATIONSHIP CARDINALITY





Cardinality refers to the uniqueness of values in a column

Ideally, all relationships in the data model should follow a
 one-to-many cardinality: one instance of each primary key,
 and many instances of each foreign key

In this example there is only **ONE instance of each Product Key** in the Product table (noted by a "**1**"), since each row contains **attributes of a single product** (name, SKU, description, price, etc.)

There are **MANY instances of each Product Key** in the Sales table (noted by an asterisk *), since there are **multiple sales for each product**

EXAMPLE: ONE-TO-ONE CARDINALITY



product_id 💌	product_name 💌	product_sku 💌	product_id 💌	product_price
4	Washington Cream Soda	64412155747	 4	\$3.64
5	Washington Diet Soda	85561191439	 5	\$2.19
7	Washington Diet Cola	20191444754	7	\$2.61
8	Washington Orange Juice	89770532250	8	\$2.59

- Connecting the two tables above using product_id creates a one-to-one relationship, since each product ID only appears once in each table
- This isn't necessarily a "bad" relationship, but you can simplify the model by merging the tables into a single, valid dimension table

product_id 🔻	product_name 💌	product_sku 💌	product_price 🔻
4	Washington Cream Soda	64412155747	\$3.64
5	Washington Diet Soda	85561191439	\$2.19
7	Washington Diet Cola	20191444754	\$2.61
8	Washington Orange Juice	89770532250	\$2.59

NOTE: this still respects the rules of normalization, since all rows are unique and capture product-specific attributes

EXAMPLE: MANY-TO-MANY CARDINALITY

Sales



Product Lookup			
product_id	¥	product_name	product_sku 🔻
4		Washington Cream Soda	64412155747
4		Washington Diet Cream Soda	81727382373
5		Washington Diet Soda	85561191439
7		Washington Diet Cola	20191444754
8		Washington Orange Juice	89770532250

date 💌	product_id 💌	transactions 💌
1/1/2017	4	12
1/2/2017	4	9
1/3/2017	4	11
1/1/2017	5	16
1/2/2017	5	19
1/1/2017	7	11

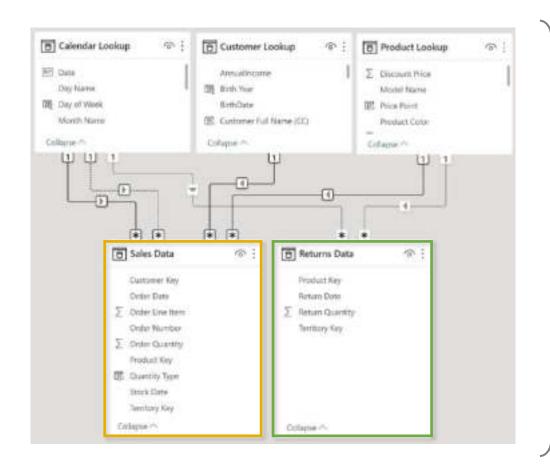


This relationship has cardinality Many-Many. This should only be used if it is expected that neither column (product_id and product_id) contains unique values, and that the significantly different behavior of Many-many relationships is understood. Learn more

- If we try to connect the tables above using product_id, we'll get a many-to-many relationship warning since there are multiple instances of product_id in both tables
- Even if we force this relationship, how would we know which product was actually sold on each date **Cream Soda** or **Diet Cream Soda**?

CONNECTING MULTIPLE FACT TABLES





This model contains two fact tables: **Sales Data** and **Returns Data**

- Since there is no primary/foreign key relationship, we can't connect them directly to each other
- But we can connect each fact table to related lookups, which allows us to filter both sales and returns data using fields from any shared lookup tables
- We can view orders and returns by product since both tables relate to Product Lookup, but we can't view returns by customer since no relationship exists

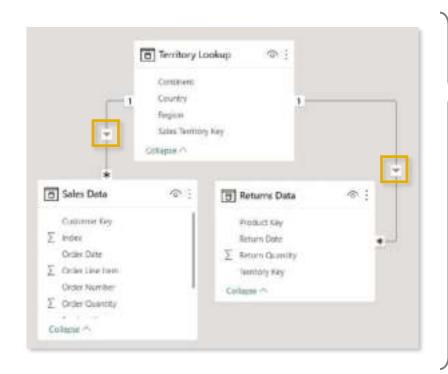


HEY THIS IS IMPORTANT!

Generally speaking, fact tables should connect through shared dimension tables, not directly to each other

FILTER CONTEXT & FLOW





Here we have two data tables (Sales Data and Returns Data), connected to Territory Lookup

The arrows show the **filter direction**, and point from the one (1) side of the relationship to the many (*) side

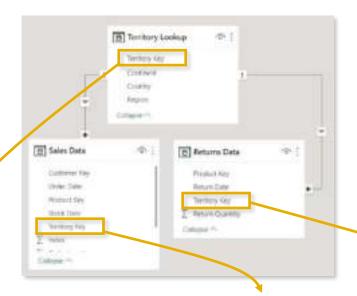
- When you filter a table, that **filter context** is passed to any related "downstream" tables, following the arrow's direction
- Filter context CANNOT flow "upstream"



PRO TIP: Arrange lookup tables above fact tables in your model as a visual reminder that filters always flow downstream

EXAMPLE: FILTER FLOW





In this model, the only way to filter both **Sales** and **Returns** data by **Territory** is to use the **Territory Key** from the lookup table, which is upstream and related to both fact tables

- Filtering using Territory Key from the Sales table yields incorrect Returns values, since the filter context can't flow to any other table
- Filtering using Territory Key from the **Returns** table yields **incorrect Sales values**, and is limited to territories that exist in the returns table



Filtering by **Territory Lookup**[Territory Key]

TerritoryKey	OrderQuantity	ReturnQuantity
1	12,513	1,828
2	40	1,828
3	30	1,828
4	17,191	1,828
5	49	28
6	894	,828
7	7,862	1,828
8	7,950	1,828
9	17,951	1,828
10	9,694	1,626
Total	84,174	1,828

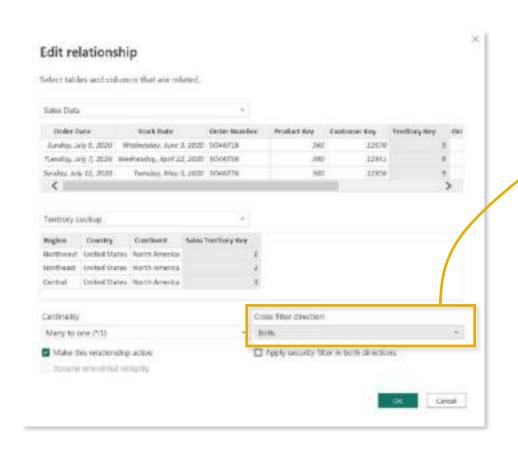
Filtering by **Sales Data**[Territory Key]

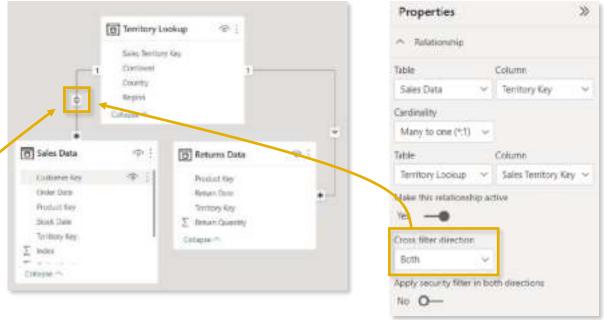
TerritoryKey	OrderQuantity	ReturnQuantity
1	84,174	270
4	84,174	362
5	84,174	1
6	174	238
7	H.174	186
8	84,174	163
9	84,174	404
10	84,174	204
Total	84,174	1,828

Filtering by **Returns Data**[Territory Key]

BI-DIRECTIONAL FILTERS





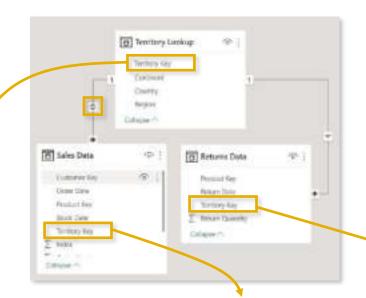


Updating the **cross-filter direction** from **Single** to **Both** allows filter context to flow in either direction

• In this example, filters applied to the **Sales** table can pass up to the **Territory Lookup** table, then down to **Returns**

EXAMPLE: BI-DIRECTIONAL FILTERS





With two-way cross-filtering enabled between **Sales** and **Territory**, we now see correct values using **Territory Key** from *either* table

- Filter context can now pass up to the **Territory Lookup** table, then downstream to **Returns**
- However, we still see incorrect values when filtering using Territory Key from the **Returns** table, since the filter context is isolated to that single table



Filtering by **Territory Lookup**[Territory Key]

TerritoryKey	OrderQuantity	ReturnQuantity
1	12,513	270
2	40	
3	30	
4	17,191	362
5	49	1
6	494	238
7	7,862	186
ñ	7,950	163
9	17,951	404
10	9,694	204
Total	84,174	1,828

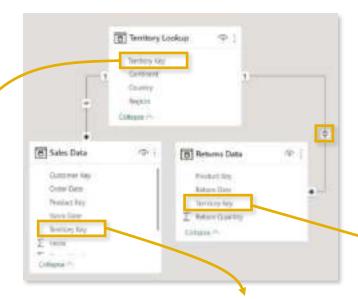
Filtering by **Sales Data**[Territory Key]

TerritoryKey	OrderQuantity	ReturnQuantity
1	84,174	270
4	84,174	362
5	84,174	1
6	174	238
7	H.174	186
8	84,174	163
9	84,174	404
10	84,174	204
Total	84,174	1,828

Filtering by **Returns Data**[Territory Key]

EXAMPLE: BI-DIRECTIONAL FILTERS





In this case, we've enabled two-way cross-filtering between the **Returns** and **Territory** tables

- As expected, we now see incorrect values when filtering using Territory Key from the Sales table, since the filter context is isolated to that single table
- While the values *appear* to be correct when filtering using Territory Key from the **Returns** table, we're **missing sales data** from any territories that didn't appear in the returns table (specifically Territories **2** & **3**)



Filtering by **Territory Lookup**[Territory Key]



Filtering by **Sales Data**[Territory Key]

TerritoryKey	OrderQuantity	ReturnQuantity
1	12,513	270
4	17,191	362
5	49	1
E	A0,894	238
7	862	86
-	7,950	163
9	17,951	404
10	9,694	204
Total	84,174	1,828

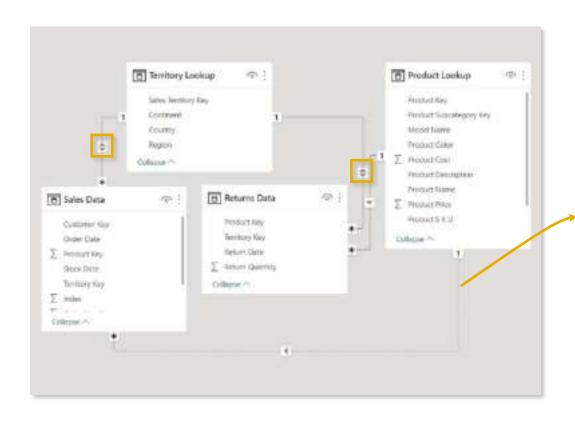
Filtering by **Returns Data**[Territory Key]

Territories 2 & 3 don't exist in the **Returns** table, so they aren't included in the filter context that passes to

Territory Lookup and **Sales**

AMBIGUITY







PRO TIP: Design your models with one-way filters and 1:many cardinality unless more complex relationships are absolutely necessary

Use two-way filters carefully, and only when necessary

• Using multiple two-way filters can cause **ambiguity** by introducing multiple filter paths between tables

You can't create a direct active relationship between Sales_Data and Product_Lookup because that would introduce ambiguity between tables Product_Lookup and Territory_Lookup. To make this relationship active, deactivate or delete one of the relationships between Product_Lookup and Territory_Lookup first.

In this example, filter context from the **Product** table can pass down to **Returns** and up to **Territory Lookup**, which would be filtered based on the Territory Keys passed from the Returns table

With an active relationship between **Product** and **Sales** as well, filter context could pass through **either the Sales or Returns table to reach the Territory Lookup table**, which could yield conflicting filter context

HIDING FIELDS





Hide in Report View makes fields inaccessible from the Report tab, but still available in **Data** and **Model** views

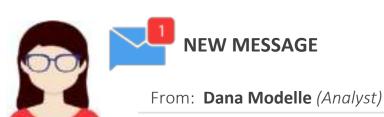
- This can be controlled by right-clicking a field in the Data or Model view, or by selecting "Is hidden" in the Properties pane
- This is commonly used to prevent users from filtering using invalid fields, reduce clutter, or to hide irrelevant metrics from view



PRO TIP: Hide the **foreign keys** in fact tables to force users to filter using **primary keys** in dimension tables

ASSIGNMENT: FILTER FLOW





Subject: Larry's gone rogue!

Hey there, we've got another problem.

Larry from Sales just sent me this screenshot. I think he must have downloaded our Power BI model and messed with some relationships, because I KNOW we had sales for product 338.

Can you help diagnose what's going on, and prevent him from doing this again?

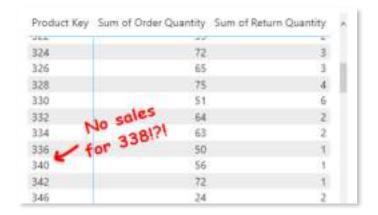
-Dana

P.S. Kevin says hi 🛣



Key Objectives

1. Replicate Larry's matrix below to diagnose what he must have done to the model*



- Which product is #338?
- Why didn't Larry's matrix show any orders?
- 2. Hide any remaining foreign keys to prevent other users from making the same mistake

*Hint: you may need to temporarily change a relationship to bi-directional

SOLUTION: FILTER FLOW





Subject: Larry's gone rogue!

Hey there, we've got another problem.

Larry from Sales just sent me this screenshot. I think he must have downloaded our Power BI model and messed with some relationships, because I KNOW we had sales for product 338.

Can you help diagnose what's going on, and prevent him from doing this again?

-Dana

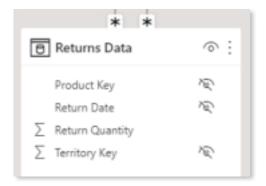
P.S. Kevin says hi 🛣



Solution Preview

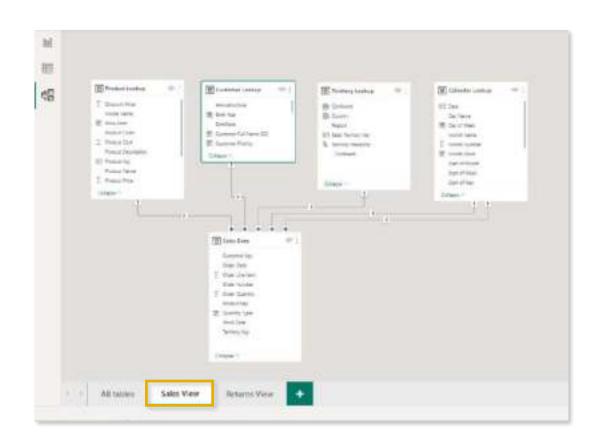
- Larry must have changed the relationship between Returns Data and Product Lookup to bi-directional, and filtered his matrix using product_id from the Returns table
 - Road bike (Road-650 Black, 44)
 - Product 338 doesn't exist in the Returns table, so it was excluded when that filter context passed to the Sales table

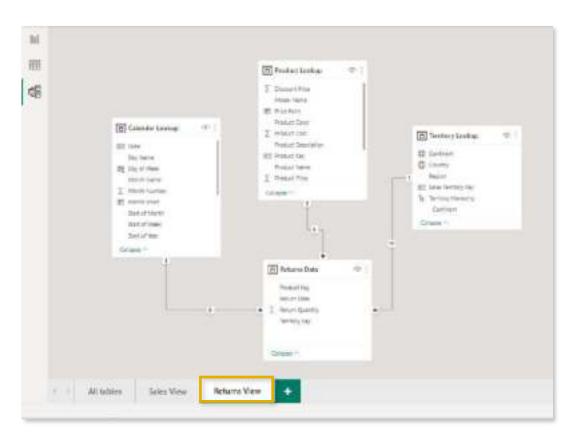
2.



PRO TIP: MODEL LAYOUTS







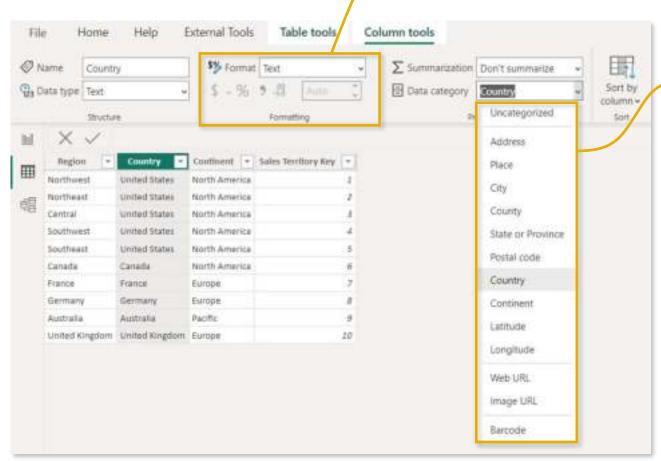
Model layouts allow you to create custom views to show specific portions of large, complex models

• Here we've created a **Sales View** displaying only tables related to sales, and a **Returns View** displaying only tables related to returns (**Note**: this doesn't actually create duplicate tables)

DATA FORMATS & CATEGORIES



Customize data formats from the Column tools menu in the
 Data view or the Properties pane in the Model view



Assign **data categories** for geospatial fields, URLs or barcodes

• This is commonly used to help Power BI map location-based fields like addresses, countries, cities, coordinates, zip codes, etc.

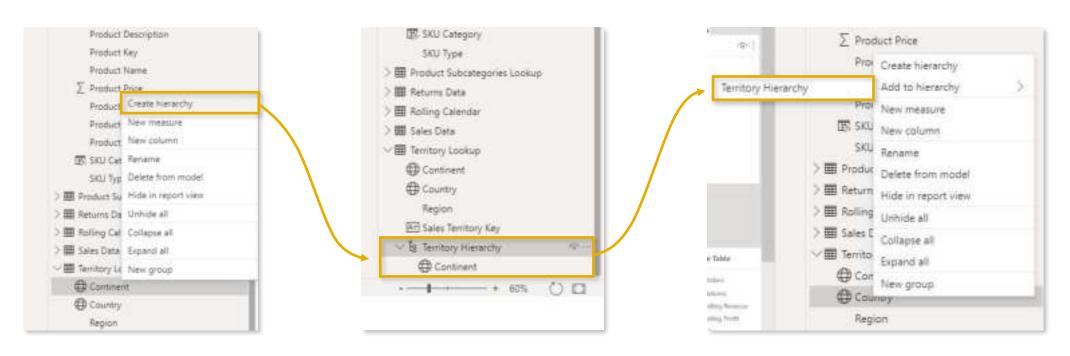


HIERARCHIES



Hierarchies are groups of columns that reflect multiple levels of granularity

- For example, a **Geography hierarchy** might include **Country**, **State** and **City** fields
- Hierarchies are treated as a **single item** in tables and reports, allowing users to "drill up" and "drill down" through each level



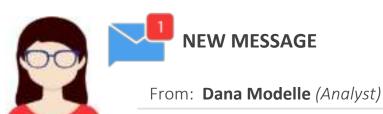
In the **Data** pane, right-click a field and select **Create hierarchy**

This hierarchy contains "Continent", and is named "Territory Hierarchy"

Right-click another field (like "Country") and select **Add to Hierarchy** (or drag it in!)

ASSIGNMENT: HIERARCHIES





Subject: Adding a date hierarchy

Good morning!

Hoping you can help with a quick request.

Since we'll be doing a lot of time-series analysis, Ethan asked us to add a date hierarchy to the model so that users can quickly view trends at any level of granularity (year, month, day, etc.)

Please get that added before our afternoon call. Thanks!

-Dana

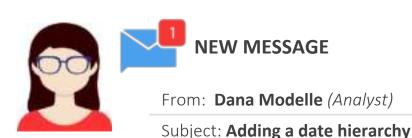


Key Objectives

- Create a new hierarchy based on the Start of Year field, and name it "Date Hierarchy"
- 2. Right-click or drag to add fields until your hierarchy contains the following (in this order):
 - Start of Year
 - Start of Month
 - Start of Week
 - Date
- Add your new hierarchy to the matrix visual (on rows) and practice drilling up and down between each level of granularity

SOLUTION: HIERARCHIES





Good morning!

Hoping you can help with a quick request.

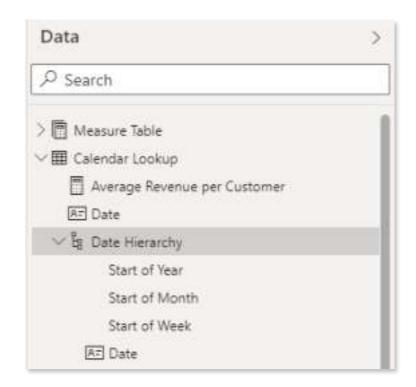
Since we'll be doing a lot of time-series analysis, Ethan asked us to add a date hierarchy to the model so that users can quickly view trends at any level of granularity (year, month, day, etc.)

Please get that added before our afternoon call. Thanks!

-Dana



Solution Preview



DATA MODEL BEST PRACTICES





Focus on building a normalized model from the start

• Leverage relationships and make sure that each table serves a clear, distinct purpose



Organize dimension tables above data tables in your model

This serves as a visual reminder that filters always flow "downstream"



Avoid complex relationships unless absolutely necessary

• Aim to use 1-to-many table relationships and one-way filters whenever possible



Hide fields from report view to prevent invalid filter context

• This forces report users to filter using primary keys from dimension tables



CALCULATED FIELDS WITH DAX



In this section we'll use **Data Analysis Expressions (DAX)** to add calculated columns & measures to our model, and introduce topics like row & filter context, iterators and more

TOPICS WE'LL COVER:

Row & Filter Context

DAX Syntax

Common Functions

Calculate

Iterators

Time Intelligence

GOALS FOR THIS SECTION:

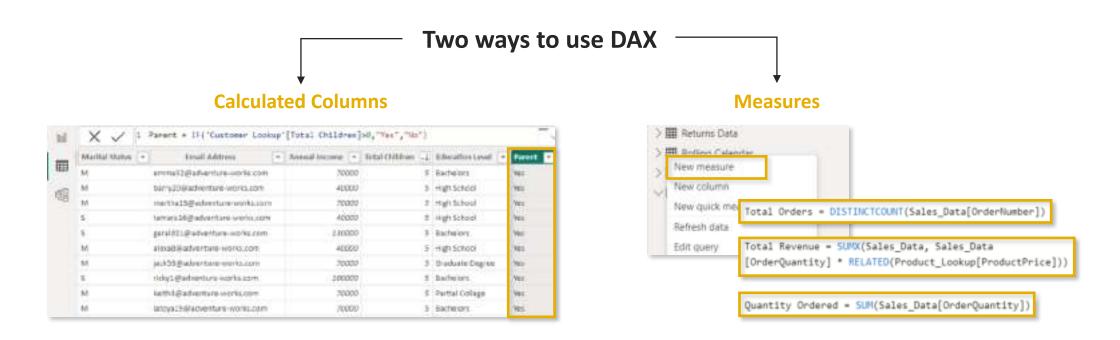
- Introduce DAX fundamentals and learn when to use calculated columns and measures
- Understand the difference between row context and filter context, and how they impact DAX calculations
- Learn DAX formula syntax, basic operators and common function categories (math, logical, text, date/time, filter, etc.)
- Explore nested functions, and more complex topics like iterators and time intelligence patterns

MEET DAX



Data Analysis Expressions (commonly known as **DAX**) is the formula language that drives the Power BI front-end. With DAX, you can:

- Go beyond the capabilities of traditional spreadsheet formulas, with powerful and flexible functions built specifically to work with relational data models
- Add calculated columns (for filtering) and measures (for aggregation) to enhance data models

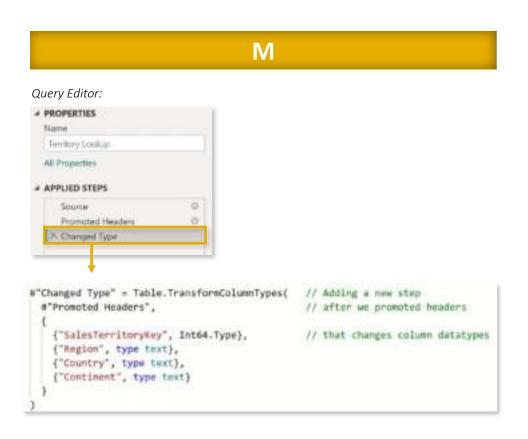


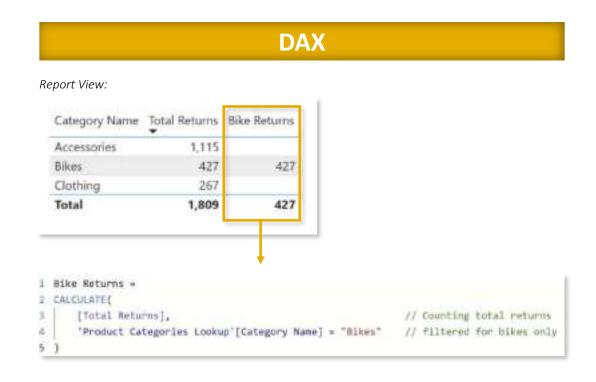
M VS. DAX



M and **DAX** are two distinct functional languages used within Power BI Desktop:

- **M** is used in the Power Query editor, and is designed specifically for extracting, transforming and loading data
- **DAX** is used in the Power BI front-end, and is designed specifically for analyzing relational data models





CALCULATED COLUMNS



Calculated columns allow you to add new, formula-based columns to tables in a model

- Calculated columns refer to **entire tables** or **columns** (no A1-style cell references)
- Calculated columns generate values for each row, which are visible within tables in the Data view
- Calculated columns understand row context; they're great for defining properties based on information in each row, but generally useless for aggregation (sum, count, etc.)

HEY THIS IS IMPORTANT!



As a rule of thumb, use calculated columns to "stamp" static, fixed values to each row in a table (or go upstream and use the Query Editor!)

DO NOT use calculated columns for aggregation – this is what **measures** are for!

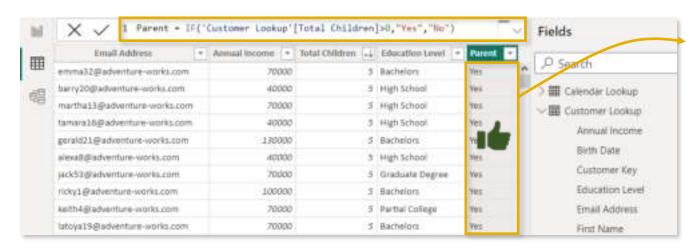


PRO TIP:

Calculated columns are typically used for filtering & grouping data, rather than creating aggregate numerical values

EXAMPLE: CALCULATED COLUMNS



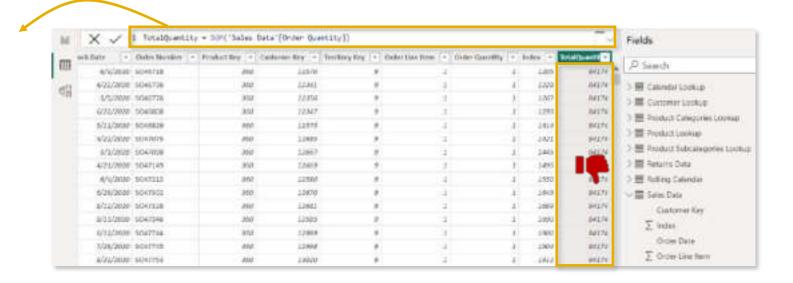


In this case we've added a **calculated column** named **Parent**, which equals "**Yes**" if the [Total Children] field is greater than 0, and "**No**" otherwise

- Since calculated columns understand row context, a new value is calculated in each row based on the value in the [Total Children] column
- This is a **valid use** of calculated columns; it creates a new row "property" that we can use to filter or segment any related data within the model

Here we're using an aggregation function (SUM) to calculate a new column named **TotalQuantity**

- Since this is an aggregation function, the same grand total is returned in every row of the table
- This is not a valid use of calculated columns; these values are statically "stamped" onto the table and can't be filtered, sliced, etc.



DAX MEASURES



Measures are DAX formulas used to generate new calculated values

- Like calculated columns, measures reference **entire tables** or **columns** (no A1-style cell references)
- Unlike calculated columns, **measures** aren't visible within tables; they can only be "seen" within a visualization like a chart or matrix (*similar to a calculated field in a PivotTable*)
- Measures evaluate based on filter context, which means they recalculate when the fields or filters around them change



HEY THIS IS IMPORTANT!

As a rule of thumb, use measures when a single row can't give you the answer, or when you need to **aggregate** values across multiple rows in a table

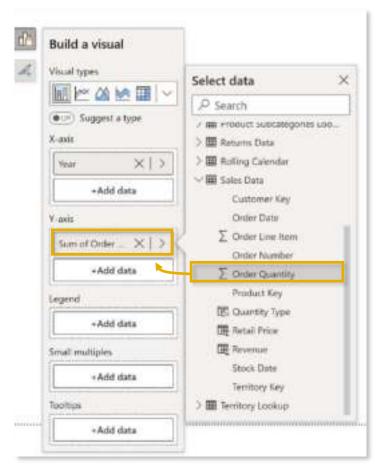


PRO TIP:

Use measures to create numerical, calculated values that can be analyzed in the "values" field of a report visual

IMPLICIT VS. EXPLICIT MEASURES





Example of an implicit measure

Implicit measures are created when you drag raw numerical fields into a report visual and manually select an aggregation mode (*Sum, Average, Min, Max, Count, etc.*)

Explicit measures are created when you actually write a DAX formula and define a new measure that can be used within the model

HEY THIS IS IMPORTANT!



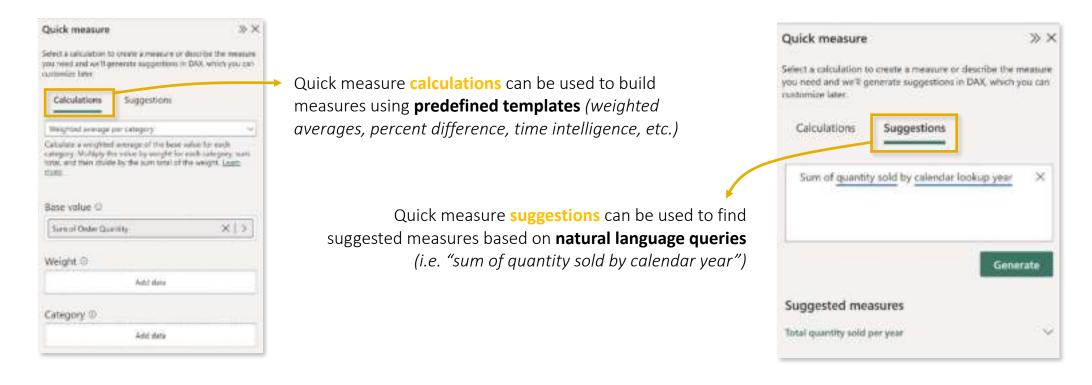
Implicit measures are only accessible within the **specific visualization** in which they were created, and cannot be referenced elsewhere

Explicit measures can be used **anywhere in the report**, and referenced by other DAX calculations to create "measure trees"

QUICK MEASURES



Quick measures automatically create formulas based on pre-built templates or natural language prompts





PRO TIP:

Quick measures can be a great learning tool for beginners or for building more complex formulas but use them with caution; **mastering DAX requires a deep understanding of the underlying theory**!